

CombiGlide 3.7

User Manual

CombiGlide User Manual Copyright © 2015 Schrödinger, LLC. All rights reserved.

While care has been taken in the preparation of this publication, Schrödinger assumes no responsibility for errors or omissions, or for damages resulting from the use of the information contained herein.

Canvas, CombiGlide, ConfGen, Epik, Glide, Impact, Jaguar, Liaison, LigPrep, Maestro, Phase, Prime, PrimeX, QikProp, QikFit, QikSim, QSite, SiteMap, Strike, and WaterMap are trademarks of Schrödinger, LLC. Schrödinger, BioLuminate, and MacroModel are registered trademarks of Schrödinger, LLC. MCPRO is a trademark of William L. Jorgensen. DESMOND is a trademark of D. E. Shaw Research, LLC. Desmond is used with the permission of D. E. Shaw Research. All rights reserved. This publication may contain the trademarks of other companies.

Schrödinger software includes software and libraries provided by third parties. For details of the copyrights, and terms and conditions associated with such included third party software, use your browser to open [third_party_legal.html](#), which is in the docs folder of your Schrödinger software installation.

This publication may refer to other third party software not included in or with Schrödinger software ("such other third party software"), and provide links to third party Web sites ("linked sites"). References to such other third party software or linked sites do not constitute an endorsement by Schrödinger, LLC or its affiliates. Use of such other third party software and linked sites may be subject to third party license agreements and fees. Schrödinger, LLC and its affiliates have no responsibility or liability, directly or indirectly, for such other third party software and linked sites, or for damage resulting from the use thereof. Any warranties that we make regarding Schrödinger products and services do not apply to such other third party software or linked sites, or to the interaction between, or interoperability of, Schrödinger products and services and such other third party software.

May 2015

Contents

Document Conventions	vii
Chapter 1: Introduction to CombiGlide	1
1.1 CombiGlide Features	1
1.2 Overview of Virtual Combinatorial Screening	2
1.3 Overview of Library Enumeration	4
1.4 Running Schrödinger Software	4
1.5 Starting Jobs from the Maestro Interface	6
1.6 Citing CombiGlide in Publications	7
Chapter 2: Enumerating a Combinatorial Library	9
2.1 Interactive Enumeration	9
2.1.1 Setting Up the Combinatorial Library	9
2.1.2 Setting Up a Fragment Collection	12
2.1.3 Creating a Fragment Collection	15
2.1.4 Enumerating the Library	17
2.2 Combinatorial Library Enumeration	17
Chapter 3: Combinatorial Screening: Defining the Chemistry	21
3.1 Obtaining Reagent and Protein Structure Files	21
3.2 Choosing the Core-Containing Molecule	22
Chapter 4: Combinatorial Screening: Preparing Structures	25
4.1 Protein Preparation and Grid Generation	25
4.2 Reagent Preparation	25
4.2.1 Choosing the Source of Reagent Structures	27
4.2.2 Selecting the Reagent Type	28
4.2.3 Setting Options for Structure Variation and Cleanup	34
4.2.4 Running the Reagent Preparation Job	35
4.2.5 Task Summary	36

4.3	Preparing the Core-Containing Molecules	36
Chapter 5:	Combinatorial Screening: The Process	37
5.1	The Combinatorial Screening Panel	37
5.2	The Combinatorial Screening Process	39
Chapter 6:	Combinatorial Screening: Defining the Core and Its Attachments	43
6.1	Selecting the Core-Containing Molecule	43
6.2	Defining Attachments	44
6.3	Saving and Loading Library Definitions	47
Chapter 7:	Combinatorial Screening: Setting Up for Docking	49
7.1	Making Glide Settings	49
7.2	Setting Glide Constraints	51
7.3	Defining the Core Poses	52
7.3.1	Selecting a Docking Mode	52
7.3.2	Specifying Structures for the Core Poses	54
Chapter 8:	Combinatorial Screening: Docking the Library	55
8.1	Running the Docking Jobs	56
8.2	Docking Results	58
Chapter 9:	Combinatorial Screening: Focusing the Library	61
9.1	Selection Strategies	61
9.2	Manual Selection of Reagents	63
9.3	Filtering	65
9.4	Results of Selection and Filtering	66
9.5	Comparing Results	69
9.6	Creating a Library	70

Chapter 10: Combinatorial Screening: Analyzing the Library	73
10.1 Chemical Feature Analysis	73
10.1.1 Features and Types	74
10.1.2 Displaying Feature and Type Analyses	75
10.1.2.1 Side Chain Features	75
10.1.2.2 Type Selectivity	76
10.1.2.3 Type Preferences	78
10.1.3 Using Chemical Feature Analysis Results	79
10.2 Active Compound Analysis	79
Chapter 11: Running Tasks from the Command Line	83
11.1 Reagent Preparation	83
11.2 Combinatorial Library Enumeration	84
11.2.1 The combgen Command	84
11.2.2 The combgen Input File	85
11.2.2.1 The Core Command	86
11.2.2.2 The Chain Command	86
11.2.2.3 The Frag Command	87
11.2.2.4 The Vlib Command	88
11.3 Library Selection	89
11.3.1 The libselector Command	89
11.3.2 The QikProp Filter File	90
11.3.3 Include and Exclude Files	90
11.3.4 Actives File	90
11.3.5 The Reagent File	91
11.3.5.1 Header Section	91
11.3.5.2 Position Sections	92
Appendix A: Creating Custom Functional Groups and Custom Minimal Capping Groups	95

Appendix B: Command-Line Tools	101
B.1 cg_add_chem_features	101
B.2 chem_features	102
B.3 cg_combine_bld	103
Getting Help	105
Glossary	109
Index	111

Document Conventions

In addition to the use of italics for names of documents, the font conventions that are used in this document are summarized in the table below.

Font	Example	Use
Sans serif	Project Table	Names of GUI features, such as panels, menus, menu items, buttons, and labels
Monospace	<code>\$SCHRODINGER/maestro</code>	File names, directory names, commands, environment variables, command input and output
Italic	<i>filename</i>	Text that the user must replace with a value
Sans serif uppercase	CTRL+H	Keyboard keys

Links to other locations in the current document or to other PDF documents are colored like this: [Document Conventions](#).

In descriptions of command syntax, the following UNIX conventions are used: braces { } enclose a choice of required items, square brackets [] enclose optional items, and the bar symbol | separates items in a list from which one item must be chosen. Lines of command syntax that wrap should be interpreted as a single command.

File name, path, and environment variable syntax is generally given with the UNIX conventions. To obtain the Windows conventions, replace the forward slash / with the backslash \ in path or directory names, and replace the \$ at the beginning of an environment variable with a % at each end. For example, `$SCHRODINGER/maestro` becomes `%SCHRODINGER%\maestro`.

Keyboard references are given in the Windows convention by default, with Mac equivalents in parentheses, for example CTRL+H (⌘H). Where Mac equivalents are not given, COMMAND should be read in place of CTRL. The convention CTRL-H is not used.

In this document, to *type* text means to type the required text in the specified location, and to *enter* text means to type the required text, then press the ENTER key.

References to literature sources are given in square brackets, like this: [10].

Introduction to CombiGlide

CombiGlide employs combinatorial technology for lead identification and optimization. It presents three distinct workflows:

1. The Interactive Enumeration workflow—Interactively construct a combinatorial library by selection of fragments for each attachment point.
2. The Library Enumeration workflow—Enumerate an entire combinatorial library from predefined reagent structure. This is most useful for relatively small libraries.
3. The Virtual Combinatorial Screening workflow—Use our proprietary technology to explore an extremely large combinatorial space in order to find side chains for a core chemical scaffold that will optimize binding to a receptor of interest.

These workflows can be used either as a prelude to combinatorial design or as a means of conventional lead discovery and optimization. Though CombiGlide carries out combinatorial chemistry in silico, the results may be used either to design focused combinatorial libraries or to evaluate a large universe of compounds for one-off or small-scale automated synthesis, or even for idea generation.

The workflows described in this manual involve substitution of the side chains on a core. The substitution of the core for a given set of side chains can also be done, with Core Hopping. This workflow is a separate product, but CombiGlide users are licensed to run the receptor-based core hopping workflow. This workflow is described in the [Core Hopping User Manual](#).

1.1 CombiGlide Features

CombiGlide includes the following capabilities and features.

Library enumeration:

- Enumerate complete combinatorial libraries.
- Untangles and minimizes structures

Interactive enumeration:

- Define and manage collections of fragments (R groups).
- Enumerate complete combinatorial libraries.

Virtual combinatorial screening:

- Performs rapid screening of large virtual combinatorial libraries against 3D targets
- Is orders of magnitude faster than docking the entire library
- Performs flexible docking using the standard and extra precision (XP) modes of Glide
- Provides multiple post-docking library selection strategies and options
- Allows for incorporation of predicted ADME properties into selection process
- Analyzes selected libraries for enrichment of actives and chemical features

All workflows:

- Provides automated reagent file preparation: 2D to 3D conversion, generation of reasonable ionization and tautomeric states, stereoexpansion, assignment of attachment points
- Uses a “core plus side chains” approach
- Offers extensive flexibility in initial core placement
- Uses an intuitive, user-friendly wizard-based GUI for setting up and monitoring jobs and for visualization of docked poses

1.2 Overview of Virtual Combinatorial Screening

CombiGlide can be used for virtual combinatorial screening, in which a focused combinatorial library is obtained at the end of the process. The steps in this process are summarized here. The last five steps constitute the combinatorial screening process (which is summarized with some more detail in [Chapter 5](#)).

1. Determine the desired synthetic approach.

The first step is to identify the chemistry associated with the desired library: the sequence of reactions that leads to the desired product. This step is discussed in more detail in [Chapter 3](#).

2. Create or obtain reagent files.

When the synthetic approach is determined, you must then select an appropriate set of reagents, and obtain files containing the structures of the reagents. This step is discussed in more detail in [Chapter 3](#).

3. Identify and prepare the core.

The core is the structural element that is constant throughout the library. The side chains are built onto the core using the reagents in the reagent files. You must supply a molecule that defines the core and ensure that it is a 3D, minimized structure, using LigPrep or MacroModel, for example. This step is discussed in more detail in [Chapter 3](#).

4. Prepare receptor and receptor grids.

The receptor must be prepared for the docking stage, and the receptor grids for the docking must be generated. These tasks are performed in the Protein Preparation Wizard and Receptor Grid Generation panels, and are described in brief in [Chapter 4](#) and in detail in the *Protein Preparation Guide* and the *Glide User Manual*.

5. Prepare reagent files.

The structures in the reagent files must be converted to 3D, all-atom structures and minimized. The bonds that will be replaced when the side chains are attached to the core must also be defined. These tasks are performed in the Reagent Preparation panel, and are described in detail in [Chapter 4](#).

6. Define attachment points and associated reagent files.

After preparing the reagent files and the core, you define the points at which the side chains from the reagents will be attached to the core to construct the library, and associate a reagent file with each attachment point. These tasks are performed in the Define Combinations step of the Combinatorial Screening panel. This step is described in detail in [Chapter 6](#).

7. Select parameters for the docking runs.

Structures are docked using CombiGlide XP docking, which makes use of Glide XP technology with some variations. The appropriate Glide options for docking can be set in the Configure Docking step of the Combinatorial Screening panel. This step is described in detail in [Chapter 7](#).

8. Define the core poses.

A set of core poses is obtained by docking selected molecules. These core poses are used as initial positions for docking the library members. In the Define Core Poses step of the Combinatorial Screening panel you can choose which molecules to dock for the core poses (they can be different from the one you selected above) and set constraints on the core position. This step is described in detail in [Chapter 7](#).

9. Dock the library members.

The docking step is done in three stages. First, the core-containing molecules are docked to generate a set of core poses. Next, structures resulting from a single substitution at any position on the core structure are docked. The unpromising reagents at each position are screened out, and the reduced list is used to dock the best of the fully substituted structures. This step is performed in the Dock Library step of the Combinatorial Screening panel, and is described in detail in [Chapter 8](#).

10. Select the optimal library.

With the docking stage complete, you can focus the library using the docking results to select the reagents that yield the best-scoring structures, and filter the results based on a selection of ADME properties. Finally, you can enumerate, and optionally dock, the focused combinatorial library. These tasks are performed in the Analyze Library step of the Combinatorial Screening panel. This step is described in detail in [Chapter 9](#).

1.3 Overview of Library Enumeration

If you already have a focused set of reagents, or if you simply want to enumerate a library, you can run a combinatorial library enumeration job. You must obtain and prepare the reagents in the same way as for library design ([Step 2](#) and [Step 5](#)); and you must prepare the core, define the attachment positions and associate the reagent files with the attachment positions ([Step 3](#) and [Step 6](#)). The attachments are defined and the library enumerated in the Combinatorial Library Enumeration panel, which is described in [Chapter 2](#).

Library enumeration can also be done interactively, in the Interactive Enumeration panel. In this process, the core molecules and the candidate side-chain molecules can be imported or sketched, the attachment positions can be marked, fragment libraries can be associated with core attachment positions, and the library can be enumerated. This process is also described in [Chapter 2](#).

1.4 Running Schrödinger Software

Schrödinger applications can be run from a graphical interface or from the command line. The software writes input and output files to a directory (folder) which is termed the *working directory*. If you run applications from the command line, the directory from which you run the application is the working directory for the job.

Linux:

To run any Schrödinger program on a Linux platform, or start a Schrödinger job on a remote host from a Linux platform, you must first set the SCHRODINGER environment variable to the installation directory for your Schrödinger software. To set this variable, enter the following command at a shell prompt:

```
csh/tcsh:      setenv SCHRODINGER installation-directory
bash/ksh:     export SCHRODINGER=installation-directory
```

Once you have set the SCHRODINGER environment variable, you can run programs and utilities with the following commands:

```
$SCHRODINGER/program &
$SCHRODINGER/utilities/utility &
```

You can start the Maestro interface with the following command:

```
$SCHRODINGER/maestro &
```

It is usually a good idea to change to the desired working directory before starting the Maestro interface. This directory then becomes the working directory.

Windows:

The primary way of running Schrödinger applications on a Windows platform is from a graphical interface. To start the Maestro interface, double-click on the Maestro icon, on a Maestro project, or on a structure file; or choose Start → All Programs → Schrodinger-2015-2 → Maestro. You do not need to make any settings before starting Maestro or running programs. The default working directory is the Schrodinger folder in your Documents folder.

If you want to run applications from the command line, you can do so in one of the shells that are provided with the installation and have the Schrödinger environment set up:

- Schrödinger Command Prompt—DOS shell.
- Schrödinger Power Shell—Windows Power Shell (if available).

You can open these shells from Start → All Programs → Schrodinger-2015-2. You do not need to include the path to a program or utility when you type the command to run it. If you want access to Unix-style utilities (such as *awk*, *grep*, and *sed*), preface the commands with *sh*, or type *sh* in either of these shells to start a Unix-style shell.

Mac:

The primary way of running Schrödinger software on a Mac is from a graphical interface. To start the Maestro interface, click its icon on the dock. If there is no Maestro icon on the dock, you can put one there by dragging it from the SchrodingerSuite2015-2 folder in your Applications folder. This folder contains icons for all the available interfaces. The default working directory is the Schrodinger folder in your Documents folder (\$HOME/Documents/Schrodinger).

Running software from the command line is similar to Linux—open a terminal window and run the program. You can also start Maestro from the command line in the same way as on Linux. The default working directory is then the directory from which you start Maestro. You do not need to set the SCHRODINGER environment variable, as this is set in your default environment on installation. To set other variables, on OS X 10.7 use the command

```
defaults write ~/.MacOSX/environment variable "value"
```

and on OS X 10.8, 10.9, and 10.10 use the command

```
launchctl setenv variable "value"
```

1.5 Starting Jobs from the Maestro Interface

To run a job from the Maestro interface, you open a panel from one of the menus (e.g. Tasks), make settings, and then submit the job to a host or a queueing system for execution. The panel settings are described in the help topics and in the user manuals. When you have finished making settings, you can use the Job toolbar to start the job.



You can start a job immediately by clicking Run. The job is run on the currently selected host with the current job settings and the job name in the Job name text box. If you want to change the job name, you can edit it in the text box before starting the job. Details of the job settings are reported in the status bar, which is below the Job toolbar.

If you want to change the job settings, such as the host on which to run the job and the number of processors to use, click the Settings button. (You can also click the arrow next to the button and choose Job Settings from the menu that is displayed.)



You can then make the settings in the Job Settings dialog box, and choose to just save the settings by clicking OK, or save the settings and start the job by clicking Run. These settings apply only to jobs that are started from the current panel.

If you want to save the input files for the job but not run it, click the **Settings** button and choose **Write**. A dialog box opens in which you can provide the job name, which is used to name the files. The files are written to the current working directory.

The **Settings** button also allows you to change the panel settings. You can choose **Read**, to read settings from an input file for the job and apply them to the panel, or you can choose **Reset Panel** to reset all the panel settings to their default values.

You can also set preferences for all jobs and how the interface interacts with the job at various stages. This is done in the **Preferences** panel, which you can open at the **Jobs** section by choosing **Preferences** from the **Settings** button menu.

Note: The items present on the **Settings** menu can vary with the application. The descriptions above cover all of the items.

The icon on the **Job Status** button shows the status of jobs for the application that belong to the current project. It starts spinning when the first job is successfully launched, and stops spinning when the last job finishes. It changes to an exclamation point if a job is not launched successfully.



Clicking the button shows a small job status window that lists the job name and status for all active jobs submitted for the application from the current project, and a summary message at the bottom. The rows are colored according to the status: yellow for submitted, green for launched, running, or finished, red for incorporated, died, or killed. You can double-click on a row to open the **Monitor** panel and monitor the job, or click the **Monitor** button to open the **Monitor** panel and close the job status window. The job status is updated while the window is open. If a job finishes while the window is open, the job remains displayed but with the new status. Click anywhere outside the window to close it.

Jobs are run under the **Job Control** facility, which manages the details of starting the job, transferring files, checking on status, and so on. For more information about this facility and how it operates, as well as details of the **Job Settings** dialog box, see the [Job Control Guide](#).

The Combinatorial Screening workflow has a different mechanism for running jobs, which is explained in the relevant places in subsequent chapters.

1.6 Citing CombiGlide in Publications

The use of this product should be acknowledged in publications as:

CombiGlide, version 3.7, Schrödinger, LLC, New York, NY, 2015.

Enumerating a Combinatorial Library

CombiGlide provides tools for simple library enumeration in addition to the tools for virtual combinatorial screening and core hopping; that is, for creating a library of ligands from a core structure to which side chains are added in user-specified positions. There are two tools for this task, the Interactive Enumeration panel, which allows you to interactively select and prepare both the core-containing molecule and the reagent molecules (fragment collections) for the library, and the Combinatorial Library Enumeration panel, which allows you to prepare the core-containing molecule but uses predefined reagent libraries to create the library. Both of these panels can be opened from the CombiGlide submenu of the Applications menu in the main window.

2.1 Interactive Enumeration

The Interactive Enumeration panel provides tools for all the tasks necessary to set up and enumerate a combinatorial library. The panel provides the means to define the attachment points on the core-containing structure and to set up the libraries for each attachment point. For each attachment point, you can also include methylene linkers between the core and the fragment that is attached. You can select structures from a predefined fragment library or create a new fragment library from a set of 3D structures. Unlike the combinatorial screening workflow, the entire combinatorial library is enumerated. Thus, this facility is not intended for use with very large combinatorial libraries.

To open the Interactive Enumeration panel, choose Applications → CombiGlide → Interactive Enumeration or Tasks → Library Design → Interactive Enumeration in the main window. From the command line, you can open this panel as an independent application with the following command:

```
$SCHRODINGER/utilities/cg_interactive_enum_dock
```

The panel has two tabs in the upper section, the Combinatorial Library tab and the Fragment Collection tab.

2.1.1 Setting Up the Combinatorial Library

To set up a combinatorial library, you need to choose a core-containing molecule, define its attachment points, and add a fragment collection at each attachment point. These tasks are done in the Combinatorial Library tab. You can use an existing library and modify it, or you can create a new library.

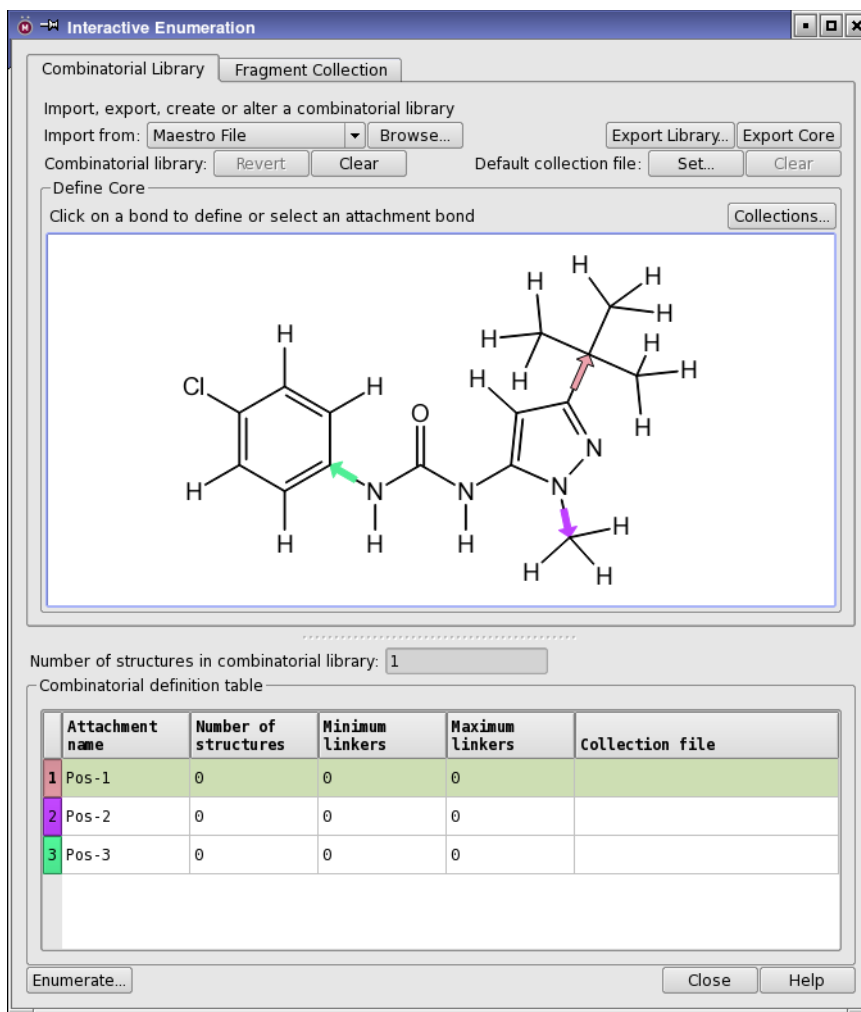


Figure 2.1. The Interactive Enumeration panel, Combinatorial Library tab.

You can import an existing library by choosing Combinatorial definition from the Import from option menu, then clicking Browse. You are prompted to save or discard changes to the current library if they are not saved, then a file selector opens, in which you can navigate to an existing library, which has the suffix `-comdef.tar.gz`. The library is loaded, the core molecule is displayed with its attachment points, and the Combinatorial definition table is filled in.

If you want to create a new library, the first step is to import a core-containing molecule. To do so, make a choice from the Import From option menu. The choices are a Maestro file (`.mae`, `.maegz`, `.mae.gz`), a combinatorial definition (`-comdef.tar.gz`), the Sketcher, the Work-

space, or the Project Table. If you choose Workspace, there must be a single entry in the Workspace. Likewise if you choose Project Table, there must be a single entry selected in the Project Table. If you choose Maestro file or Combinatorial definition, a file selector opens, in which you can navigate to and select the file that contains the structure you want to use to define the core. The first structure in a Maestro file is the structure that is used. If you choose Sketcher, click Sketch to open the 2D Workspace panel, and you can draw the structure in this panel. When you click OK, the structure is converted to 3D with LigPrep and imported. For information on this panel, see [Section 5.5](#) of the *Maestro User Manual*.

When you have imported the core-containing molecule, it is displayed in the display area with any existing attachment bonds, or core positions, marked by colored arrows. The arrow color matches the rows in the Combinatorial definition table. The head of the arrow indicates the fragment or atom that will be removed when a fragment (R group) is attached in this position.

You can then pick bonds in the structure to define the attachment points, or core positions. If you right-click on an existing core position, the shortcut menu allows you to switch the direction of the arrow or delete it. Clicking again on the bond also changes the direction of the arrow. The direction is only changed, however, if it would not invalidate existing core positions. Each time you create a new core position, it is added to the Combinatorial definition table. The selected position is highlighted in olive-green in the table and the arrow is outlined in black in the display area.

To delete a core position, right-click on the arrow in the display area or in the table row and choose Delete.

The Combinatorial definition table lists each core position, along with information about the position and its attachments: the number of structures in the fragment collection used at this position, the minimum and maximum number of methylene linkers, and the name of the fragment collection file. The fragment collection itself for a given core position is shown in the Fragment Collection tab, which is described in the next section.

When a core position is first selected, a default name is created and the numbers in the columns are set to zero. You can edit the name by clicking in the table cell and changing the text. The fragment collection for this position in the Fragment Collection tab is initially empty, by default. To associate a fragment collection with this position, you can create one in the Fragment Collection tab, or you can read one from file. You can also use the Default collection file buttons to set up or clear a default fragment collection, which is automatically added to a position when you create it. Click Set, and navigate to the collection file in the file selector that opens, to define a default collection; click Clear to return to having no default collection.

To read a new collection from file, right-click on the selected row of the table and choose either Replace Collection or Append to Collection. It does not matter which you choose because the existing collection is empty. In the file selector that opens, you can navigate to and select a

fragment collection, which has an extension of .bld. The directory selection includes Schrödinger Collections, which supplies a default set of fragments. The directory selection also includes My Collections, which is set to the `cg/interactive_fragments` subdirectory of your user resources directory (`$HOME/.schrodinger` on Linux and Mac, `%APPDATA%\Schrodinger` on Windows) as a convenient location for storing your fragment collections. Once the collection is read in, it is shown in the Fragment Collection tab. The number of fragments in the file is displayed in the Number of structures column.

If you have created a fragment collection in the Fragment Collection tab, you can apply it to the current core position by right-clicking in the table and choosing **Apply Fragment Collection**, which replaces the current collection with the entire collection, or choosing **Apply Selected Fragments**, which replaces the current collection with the fragments that are selected in the Fragment Collection tab. A default name for the fragment file is created. You can change it by right-clicking on the table row and choosing **Rename**.

Another way of adding a collection is from the Collections panel, which you open by clicking **Collections**. You can then choose a collection from the list and drag it onto the desired attachment bond of the 2D structure in the Define Core section.

As well as the fragments that you define, the combinatorial library can include methylene linkers between the core and the fragment. The number of linkers is determined by the values given in the Minimum linkers and Maximum linkers columns, which you can edit to set the desired values.

If you changed a library definition and want to revert to the original definition, click **Revert**. All changes since the last time the library was saved are discarded and the library is reloaded.

To start a new library definition, you can click **Clear** to clear the current library definition.

When you have a library that you want to save, click **Export Library**. The library must have a fragment collection associated with each core position. A file selector opens, in which you can navigate to the desired location and name and save the library definition.

2.1.2 Setting Up a Fragment Collection

The Fragment Collection tab provides the tools to design your own fragment collection. You can import fragments, create fragments, delete fragments, and export fragments. You can either set up a fragment collection independent of a core position, or you can set up a fragment collection for a particular core position. If a core position already has a collection file associated with it, the contents of this file are loaded as the current collection when you select the core position.

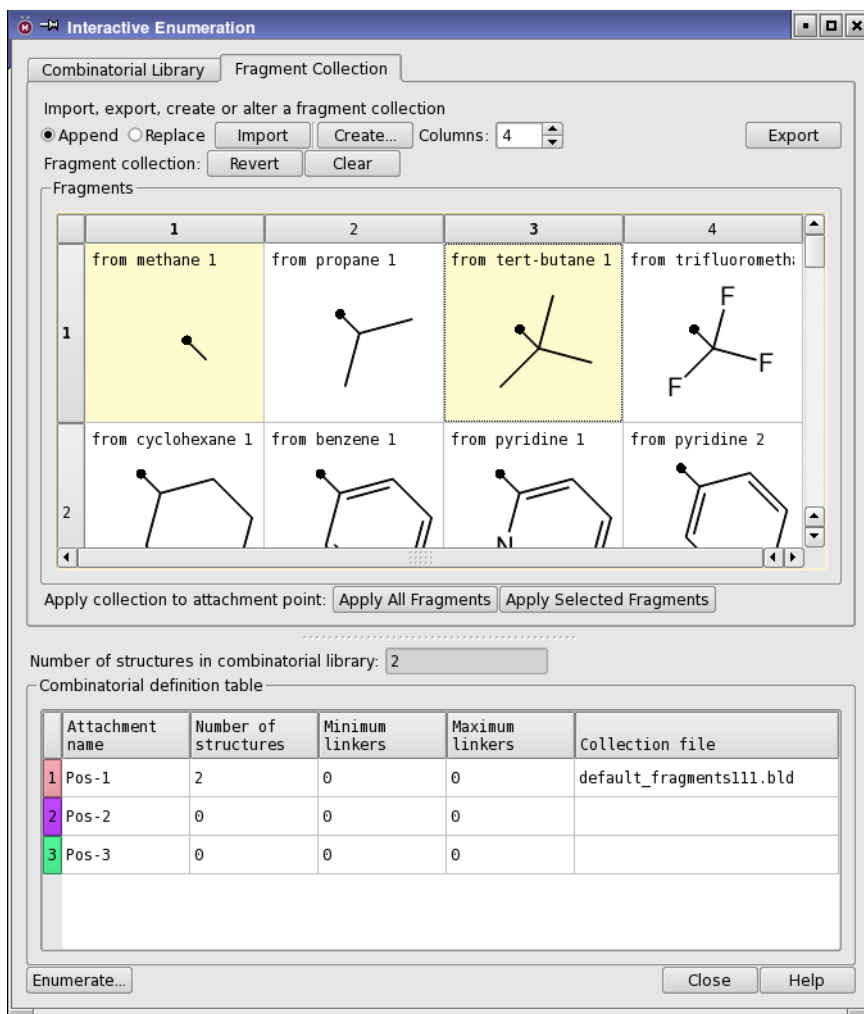


Figure 2.2. The Interactive Enumeration panel, Fragment Collection tab.

To add fragments to a collection, you can import them or create them. Creating new fragments is done by clicking Create and working in the Create Fragment Collection panel. This task is described in the next section. In addition to using the Create Fragment Collection panel, you can prepare fragments in the Reagent Preparation panel and then import them—see [Section 4.2 on page 25](#) for more information.

You can import fragments from a file, from the Project Table, or from the Workspace. The fragments must be already prepared with the information on the attachment bond, otherwise you will not be able to import them. If you import from a file, the file must have a .bld extension.

sion. If you import an entry group from the Project Table, the collection is named after the group name. Before you click **Import** to import the structures, you should decide whether you want to add the structures to the existing collection, by selecting **Append**, or replace the existing collection with the new structures, by selecting **Replace**.

The 2D structures of the current fragment collection are displayed in the center of the tab. The filled circle indicates the connection point to the core. You can set the number of columns of structures in the display area in the **Columns** box. You can select multiple structures by dragging or with shift-click and control-click. When you right-click on a structure, the shortcut menu offers the choice of deleting this structure or the selected structures, or retaining only this structure or only the selected structures.

If you are editing the fragment collection for a particular core position, the collection that is used for that position is not changed unless you specifically change it. This means that you can experiment with the collection before you commit to applying it to the core position. When you want to apply the changes, you can choose one of the following options from the shortcut menu of the Combinatorial definition table:

- **Apply Fragment Collection**—replace the collection file with the entire fragment collection
- **Apply Selected Fragments**—replace the collection file with the fragments that are selected in the fragment display area.

You can also use one of the **Apply** collection to attachment point buttons: **Apply All Fragments** applies all fragments to the current attachment point, and **Apply Selected Fragments** applies only the selected fragments to the current attachment point. The table row for the core position is colored red when you edit the collection, to indicate that the displayed collection is not the same as the collection associated with the position.

If you want to discard the changes you have made to an existing collection, click **Revert**. The changes are discarded and the structures are read in again from the file. If you want to clear all the fragments from the current collection, click **Clear**. In both cases you are prompted to confirm the action.

When you want to save a collection, you can click **Export** to export it to a file or to the Project Table. A menu is displayed so that you can choose the destination. If you export to the Project Table, it is exported as an entry group named after the collection file. If you export to a file, the default file location is **My Collections**, which is described in the previous section.

2.1.3 Creating a Fragment Collection

If you want to create a new fragment collection by selecting from a set of structures, you can do so in the Create Fragment Collection panel. In this panel, you can import or draw structures, select the ones that you want, define their attachment bonds, and add them to the current fragment collection in the Interactive Enumeration panel.

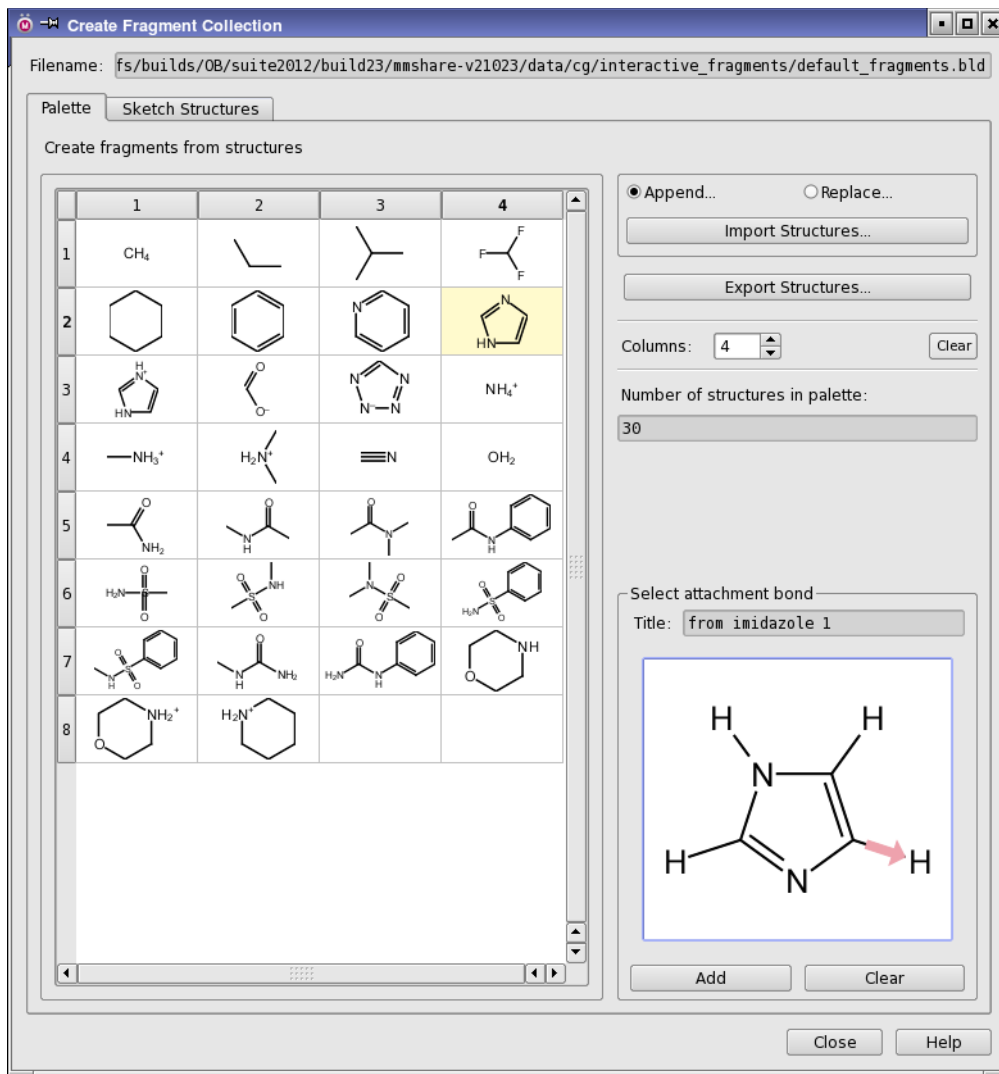


Figure 2.3. The Create Fragment Collection panel.

To open the Create Fragment Collection panel, click Create in the Interactive Enumeration panel. The panel opens with the Palette tab displayed.

The palette—the display area for the structures—is initially empty. The palette is a table of structures that can be selected for inclusion in the fragment collection, with an attachment bond defined. You can set the number of columns in the Columns box, and resize the table cells by dragging the row or column boundary in the table margin. You can clear the palette by clicking Clear.

To import structures into the palette, click Import Structures. A menu is displayed that allows you to choose whether to import from a file, the selected entries in the Project Table, or the Workspace. If you choose Import from File, a file selector opens, and you can navigate to and select the file. The file can be a Maestro file, without attachment bonds defined, or a reagent file with attachment bonds defined (.bld). The file selector opens in a directory that contains a default set of structures that are useful in a palette. When you click Open, the file name is displayed in the Filename text box, and the structures are shown in 2D form in the palette.

If you already have structures, you can append to the list by selecting Append and importing structures. If you want to replace the existing structures with a new set, select Replace before importing the structures.

To draw structures, click the Sketch Structures tab. In this tab you can draw structures with the 2D Sketcher, then add them to the palette. You can place multiple structures on the drawing area, and when you click Add Structures, all of them are added as independent fragments, after checking for duplicates. For information on the 2D Sketcher, see [Section 5.5](#) of the *Maestro User Manual*.

You can select multiple structures in the palette (with shift-click and control-click). The shortcut menu offers the choice of keeping only the selected structures or deleting them, or keeping only the structure clicked on or deleting it. This allows you to define new palette collections for later use.

If you want to save a particular set of structures, click Export Structures. The menu that is displayed offers the choice of exporting to a Maestro file or to the Project Table. Attachment bond information is not written with the structures. If you choose to export to a file, a file selector opens, in which you can navigate to a location and name the file. The file selector opens in a location beneath your home directory that can be used as the default location of any structure files that you might like to save for reuse in a palette. The export option is useful for creating a new file for future use in a palette after importing and possibly deleting structures from several files.

Once you have structures in the palette, you can proceed to select a structure, add an attachment bond, and then add the structure to the current fragment collection in the Interactive Enumeration panel.

When you click on a structure in the palette, the structure is displayed in the Select attachment bond section, so that you can define the attachment bond. The stereochemistry is indicated by wedge and dashed bonds. The title of the structure is displayed in the Title text box, and the structure itself in the display area. To select the attachment bond, click on a bond in the display area. Clicking a second time on the same bond changes the direction of the attachment bond, if the bond is not terminal. The arrow on the bond points towards the part that is removed to attach the fragment to the core.

To add this fragment to the collection, click Add. The fragment flashes, and it is added to the current Fragment Collection display of the Interactive Enumeration panel. It is therefore a good idea to have both panels visible and next to each other, so that you can see the fragment being added to the fragment collection.

If you want to add the same fragment with a different attachment point, you can select a different attachment bond and click Add. The fragment is added to the collection again with the new attachment bond.

2.1.4 Enumerating the Library

When you have defined a combinatorial library, you can then proceed to enumerate it. When you click Enumerate, the Start dialog box opens, and you can set job options and start the enumeration job.

The results can be appended to the Project Table or written to the current directory. The default is to append the results to the Project Table.

2.2 Combinatorial Library Enumeration

The combinatorial library enumeration is carried out by substitution of side chains on the core structure with fragments from the reagent structures. For both the core and the reagents, the bonds to be broken must be designated. The fragments from the reagents are then attached to the core at the site of the broken bond. This process can be regarded as "growing" the new side chain onto the core in place of the old side chain, and the bond that is replaced (with its attached atoms) is called the grow bond. These bonds are not necessarily the bonds that are broken and formed in the real chemical reaction, but represent a means of varying the side chains on a core structure that might include part of the real reagent.

To generate a combinatorial library, you must have a 3D, minimized structure for the core-containing molecule whose side chains you want to replace, and a set of prepared reagent files that define the replacements. For information on preparing the core structure and the reagent files, see [Chapter 4](#).

The Combinatorial Library Enumeration panel is essentially the same as the Define Combinations step of the Combinatorial Screening panel, which is described in [Chapter 6](#). The procedure is summarized here.

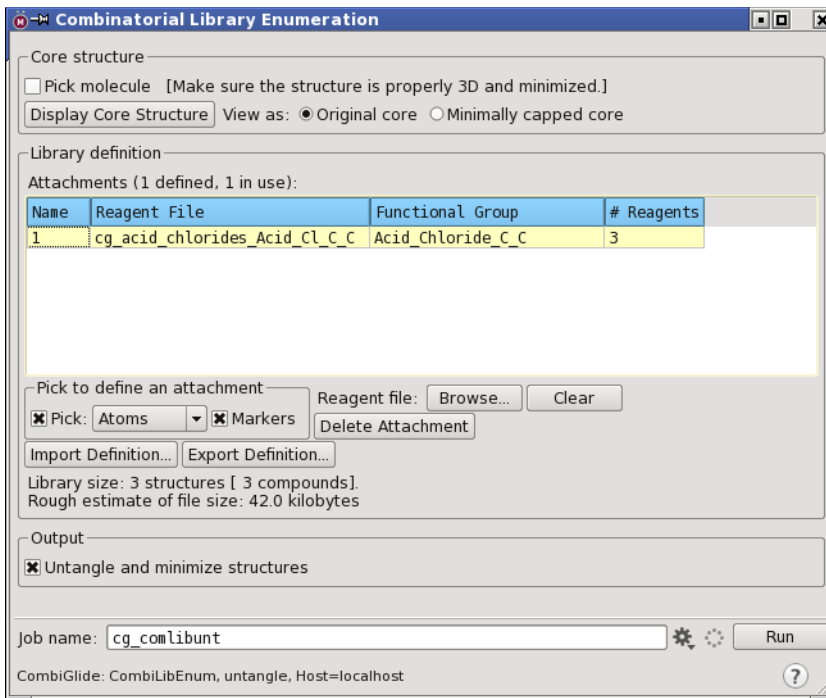


Figure 2.4. The Combinatorial Library Enumeration panel.

When you first open the Combinatorial Library Enumeration panel, Pick molecule is selected in the Core structure section, and most other controls are unavailable. After you have selected a core-containing molecule, the other controls become available.

To define the core-containing structure:

1. Ensure that the structure is displayed in the Workspace.
2. Click on an atom in the core-containing molecule.

The core-containing molecule is now defined, and you can specify the attachments. Pick molecule is deselected, and Pick is selected in the Pick to define an attachment section.

To define an attachment:

1. Pick two atoms that define the position of the attachment, the atom that is kept first, then the atom that is deleted. (The order is important.)

A turquoise arrow is shown over the bond, and the Select Reagent File dialog box opens.

2. Select the appropriate reagent file, and click OK.

The reagent files have a .bld extension: they are Maestro-format files that have special information for the Maestro builder, which does the library enumeration. The dialog box closes, and details of the attachment are listed in the Attachments table. The arrowhead changes to a sphere and the color changes to gold. The size of the library with the current set of attachments is given in the lower portion of the panel.

If you choose the wrong location for an attachment position, you can delete it by selecting it in the table and clicking **Delete Attachment**. If you attach the wrong file, you can select the attachment position in the table, and click **Browse** (below the table) to select the correct file.

When you have defined an attachment, the size of the resultant library is displayed at the foot of the Library definition section.

It is recommended that you leave the **Untangle and minimize structures** option selected. Untangling ensures that (for example) any chains that pass through the middle of rings during the build process are removed from the rings and the structures are properly minimized. However, if you are not interested in having accurate structures, you can deselect this option.

To start the library enumeration job, do one of the following:

- Click **Run** to run the job immediately with the default job settings.
- Choose **Job Settings** from the **Settings** button menu, make job settings, and click **Run**.

The host you choose should have temporary storage space of about 3 to 4 times the size of the library. This space is required for the untangling and minimization of the library members.

The job takes about 45-60 seconds per output structure when each structure is minimized and untangled. The library is written to a Maestro file named *jobname.mae*. You can specify the directory that this file is written to in the **Job Settings** dialog box.

If you choose to untangle the structures, you can distribute the job that performs the untangling over multiple processors. In the **Job Settings** dialog box, specify the number of subjobs in the **Separate untangling into N subjobs** text box, then select an option for distributing the subjobs. You can distribute the subjob over the maximum available processors on the selected host, or specify the number of processors to use. For optimum load balancing, the number of subjobs should be several times the number of processors used.

Combinatorial Screening: Defining the Chemistry

The first step in the design of a focused library using CombiGlide is to identify the chemistry associated with desired library. Most likely, this will be done in collaboration with one of your medicinal or combinatorial chemists.

For example, if you were interested in generating a focused pyrazole library in an attempt to identify inhibitors of p38 MAP kinase, your synthetic chemists might suggest the approach described in [Figure 3.1](#).




Figure 3.1. Synthetic route to pyrazole library.

3.1 Obtaining Reagent and Protein Structure Files

Once the synthetic approach has been determined, it will be apparent which types of reagents will be required for the library synthesis. For example, in the pyrazole library described in [Figure 3.1](#), the library will be generated by the combination of acid chlorides, hydrazines, and isocyanates.

You must then obtain files that contain the structures of the reagents that you want to use. For each reagent, the structures must be contained in a single, multistructure file. You can provide

files that contain only the structures for particular reagents, or you can provide files that contain structures for a variety of reagents, from which CombiGlide will select out the specified types of reagents. In the pyrazole library example, you could supply one file containing all of the acid chlorides to be considered, one file for the hydrazines, and one for the isocyanates; or you could supply one large file containing structures for all three reagent types. The reagent files must be in either SD or Maestro format.

It is highly recommended that the reagent files contain a set of reagents that has been prefiltered to remove redundant compounds and any reagents that are not compatible with the synthetic approach. Prefiltering will reduce the amount of computational work needed to obtain a focused library and is more likely to lead to useful results.

If you do not already have a library of reagents, you can find links to sources of compounds at the HTScreening.net web site, <http://www.htscreening.net/home>, or the ZINC web site, <http://blaster.docking.org/zinc>. Most vendors supply structures of the compounds in 2D SD format.

You will also need to obtain a file containing the structure of the receptor. Protein structures can be obtained from the Protein Data Bank (PDB), <http://www.rcsb.org/pdb>. You can import a PDB structure using Project → Get PDB, or by entering the PDB ID in the Protein Preparation Wizard panel (which you will need to use to prepare the receptor). Both of these are accessible from the toolbar.



If you have the PDB database installed (and the environment variable `SCHRODINGER_PDB` is set to identify the database), the protein structure is obtained from the database directly, otherwise it is retrieved from the web site. See [Section 3.1](#) of the *Maestro User Manual* for more information.

3.2 Choosing the Core-Containing Molecule

CombiGlide works by building side chains onto defined positions on a core-containing structure. You must therefore supply a structure that contains the core in addition to supplying the reagents. The reagent structures are used as a source of side chains. Even though in the actual chemical reactions, atoms from the reagents can be included into the core, these atoms are discarded when the side chains are added to the core.

There are some limitations on how you can select the reagents and attachment positions, which may affect the choices you make for the core-containing molecule.

- The bond that is replaced in the core-containing molecule and in the reagent when the side chain is added must be a single bond. In the pyrazole example in [Figure 3.1](#), the bond

that is replaced in the isocyanate is the single C–R' bond. The C=N bond in the isocyanate could not be chosen because it is a double bond, even though the corresponding bond in the core-containing molecule is a single bond.

- The bond must not be in a ring. The hydrazine N–N bond in the pyrazole example could not be chosen because it forms part of the pyrazole ring.

Core-containing molecules are also docked to define core poses, which are used to build and dock the substituted structures. By default, the minimally capped core molecule is used, but you can choose other molecules to dock for the core poses. When you choose core-containing molecules for use in defining core poses, it is often advantageous to select ones that are not so small that they can dock in poses that are unreasonable for the final substituted molecules. It does not matter how long or short the side chains are in the core-containing molecules, since they will be replaced. A good choice would be a known active.

Combinatorial Screening: Preparing Structures

When you have decided on the biological target and the synthetic route to the library you want to generate, and obtained structure files for the reagents and the receptor, you can proceed to the preparation of the structures for CombiGlide. Both the protein structure and the reagent structures must be properly prepared. You can perform these tasks in the Protein Preparation Wizard and Reagent Preparation panels. Following protein preparation, you must set up grids for Glide docking, which you can do in the Receptor Grid Generation panel.

4.1 Protein Preparation and Grid Generation

A prepared protein structure is required for CombiGlide docking. Protein preparation can include simplifying multimeric complexes, deleting unwanted waters or cofactors, correcting serious errors in the protein structure, adjusting bond orders, ionization states, and formal charges of the protein, metal ions, cofactors, and ligand, adding hydrogens, neutralizing appropriate amino acid chains, reorienting side-chain hydroxyl and thiol groups, and relieving steric clashes. The full procedure for protein preparation is described in [Section 1.1](#) of the *Protein Preparation Guide*. The entire process can be performed in the Protein Preparation Wizard panel, which is described in [Chapter 2](#) of the *Protein Preparation Guide*. You can open the Protein Preparation Wizard panel from the Workflows menu.

After running the protein preparation job, it is always advisable to check the structure for any remaining problems.

When you have prepared the protein, you must generate grids for docking with CombiGlide, using the Receptor Grid Generation panel. The grid generation process is described in detail in [Chapter 4](#) of the *Glide User Manual*. You should refer to this chapter and follow the procedures in it to generate your grids. You can open the Receptor Grid Generation panel from the CombiGlide submenu of the Applications menu. This panel is also identical to the Glide version of the panel.

4.2 Reagent Preparation

Reagent preparation ensures that the input structures are all-atom, 3D structures, and that they have the appropriate information stored with them to construct the molecules that are docked or used for library enumeration. The main tasks in the reagent preparation process are to select the source of reagent structures, select a reagent type (a functional group), identify the bond in

the functional group that is replaced when the reagent is added to the core, and perform 2D-to-3D conversion, structure variation and cleanup. The first three of these tasks are performed in the Reagent Preparation panel, and options for the fourth are set up in this panel. A job is then run to obtain the structures from the input file and prepare them for use with CombiGlide.

To open the panel, choose Applications → CombiGlide → Reagent Preparation. This panel contains three sections: one for selecting the input structure file, one for selecting the functional group, and one for setting options for structure conversion, variation, and cleanup.

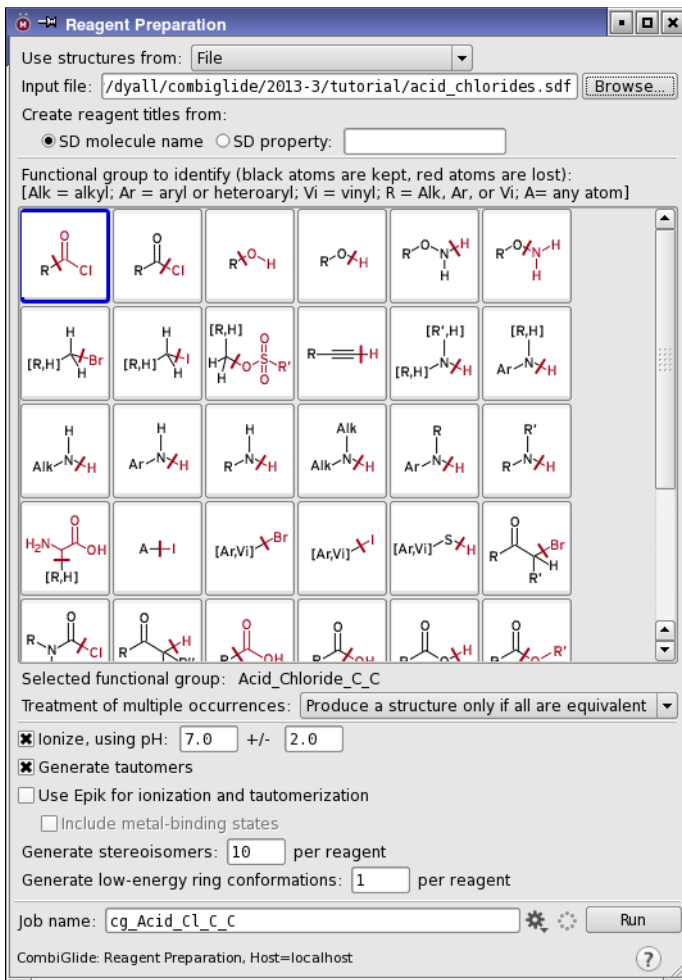


Figure 4.1. The Reagent Preparation panel.

You must run a reagent preparation job for each reagent type that will be used in combinatorial screening or in combinatorial library enumeration. A reagent type corresponds more or less to a functional group, plus information on which bond is to be replaced in the build process. The tasks described briefly above and in detail below must be performed for each reagent type that you plan to use. In the pyrazole example from [Chapter 3](#), there are three reagent types: acid chlorides, hydrazines, and isocyanates.

A reagent preparation job is a special kind of LigPrep job that performs some tasks needed by CombiGlide in addition to the normal LigPrep tasks of ligand preparation and variation. For more information on LigPrep, see the [LigPrep User Manual](#).

4.2.1 Choosing the Source of Reagent Structures

The first task is to choose the source of the reagent structures for a given reagent type. Because the reagent preparation job matches the selected functional group in each structure to a SMARTS pattern, structures that do not match are filtered out. The structures that match are written to a file and passed on to the next stage. The structure input file can therefore contain any structures, not just those that match the desired functional group. However, it is highly recommended that the input file you choose is prefiltered to eliminate redundant structures and structures that are not compatible with the synthetic approach.

There are some restrictions on the structures that you can use as input. If a structure contains more than one instance of the selected functional group, it will not be written to the output file unless it is symmetrical, in which case only one of the bonds will be designated as the bond to be broken in the process of building the library member from the reagents.

Once you have opened the Reagent Preparation panel, you can choose the source of structures from the Use structures from option menu. You can select structures from the Workspace, the Project Table, or from an external file. Structures from the Workspace or the Project Table are written to a Maestro file that is passed to the reagent preparation job.

If you choose File from the Use structures from option menu, you can enter the path to the file in the Input file text box, or click Browse to navigate to the file. The file must be in Maestro (.mae), SD (.sd or .sdf), or SMILES (.smi) format. The structures can be 2D or 3D, and are converted to all-atom 3D structures by the reagent preparation job.

If you selected an SD file, you can also specify the source of the reagent titles. Reagent titles are used for identification in CombiGlide, so it is important to have unique titles. If you do not make them unique, CombiGlide will do so for you. You can select SD molecule name to use the molecule name from the SD file, or select SD property and enter a property name in the text box to use the property for the reagent titles.

4.2.2 Selecting the Reagent Type

Selecting the reagent type involves choosing a functional group along with the bond that is to be replaced when the reagent is used to build structures for the library. A variety of predefined functional groups with corresponding bonds to be broken are displayed as a set of buttons in the center of the Reagent Preparation panel. The icon for each functional group marks the bond that is broken to attach the reagent to the core with a red line across the bond. The part of the reagent that is discarded is shown in red.

To select a reagent type, click the button that corresponds to the desired combination of a functional group and a bond to break. The selection is displayed below the scrolling region. The “short name” and the “long name” of the selected functional group are displayed below the icon area. Both names consist of a part that identifies the type of compound (acid chloride, primary amine, and so on) and a part that identifies the bond that is broken, separated by an underscore character. The part that identifies the bond consists of two element names separated by an underscore. The first name defines the atom that is kept when the reagent is added to the core, and the second name defines the atom that is discarded, along with anything attached to it. For example, `Acid_Chloride_C_C` is the long name for acid chlorides in which the carbon-carbon bond of the acid chloride group is broken. In this case the entire -COCl group is discarded, and the R group is added to the core.

A single class of compounds can have multiple icons, corresponding to different bonds that are broken. You should make sure that you have selected the icon with the bond to be broken that is the most appropriate for your molecular system. For example, acid chlorides can have the C–C bond or the C–Cl bond broken.

The chosen functional group might be matched multiple times in a reagent molecule, and a decision must be made about how to handle these occurrences of the functional group. Three choices are offered on the Treatment of multiple occurrences option menu:

- Produce all variants—Prepare the reagent molecule by annotating each occurrence of the functional group as a separate structure.
- Ignore the molecule—Discard the reagent molecule. This choice produces a reagent file in which there is only one occurrence of the functional group in each structure.
- Produce a structure only if all are equivalent—If the occurrences of the functional group are all equivalent (as for example the acid group in oxalic acid), annotate the functional group for one of the occurrences, otherwise discard the molecule.

A detailed description of each functional group is provided in [Table 4.1](#). The functional groups are defined using SMARTS patterns.

If the functional group that you want to use is not listed, you can add a custom functional group. Details on the requirements and procedure are given in [Appendix A](#).

Table 4.1. Predefined functional groups with definitions

Structure	Short name Long name	Definition of R, R', R'', Alk, Ar, Vi, A
	Acid_Cl_C_C Acid_Chloride_C_C	R can be anything with a carbon attached to the carbonyl.
	Acid_Cl_C_Cl Acid_Chloride_C_Cl	See above
	Alc_C_O Alcohol_C_O	R can be an alkyl or aryl group. R cannot have a carbonyl attached to the oxygen of the alcohol.
	Alc_O_H Alcohol_O_H	See above
	Alkoxyamine_N_H Alkoxyamine_N_H	R can be H or anything with a carbon attached to the oxygen of the alkoxyamine.
	Alkoxyamine_O_N Alkoxyamine_O_N	See above
	Alk_Br_C_Br Alkyl_Bromide_C_Br	R can be almost anything: H, alkyl, aryl, alkenyl, alkynyl; alkoxy, aryloxy, alkoxy carbonyl, or aryloxy carbonyl with oxygen attached to CH ₂ ; silyl with silicon attached to CH ₂ ; alkylamino, arylamino, alkylaminocarbonyl, or arylaminocarbonyl with nitrogen attached to CH ₂ ; alkylthio, arylthio, alkylsulfinyl, arylsulfinyl, alkylsulfonyl, or arylsulfonyl with sulfur attached to CH ₂ ; ketone with carbonyl attached to CH ₂ ; cyano. R cannot be chloro, iodo, chlorocarbonyl (carbonyl attached to CH ₂).

Table 4.1. Predefined functional groups with definitions









Structure	Short name Long name	Definition of R, R', R'', Alk, Ar, Vi, A
	Alk_I_C_I Alkyl_Iodide_C_I	Same as for alkyl bromides
	Alk_Sulf_C_O Alkyl_Sulfonate_C_O	Same as for alkyl bromides
	Any_I_A_I Any_Iodide_A_I	A can be any atom for which force fields are available.
	Alkyne_C_H Alkyne_C_H	R can be H, alkyl, aryl, silyl.
	Amine_Gen_N_H Amine_General_N_H	R can be H, alkyl, aryl. R cannot have a carbonyl carbon attached to the nitrogen of the amine. R' can be H, alkyl, aryl. R' cannot have a carbonyl carbon attached to the nitrogen of the amine.
	Amine_Gen_Ar_N_H Amine_General_Aryl_N_H	Ar can be aryl. R can be H, alkyl, aryl. R cannot have a carbonyl attached to the nitrogen of the amine.
	Amine_Prim_Alk_N_H Amine_Primary_Alkyl_N_H	R can be H, alkyl. R cannot have a carbonyl attached to the nitrogen of the amine.
	Amine_Prim_Ar_N_H Amine_Primary_Aryl_N_H	Ar can be aryl.

Table 4.1. Predefined functional groups with definitions









Structure	Short name Long name	Definition of R, R', R'', Alk, Ar, Vi, A
	Amine_Prim_Gen_N_H Amine_Primary_General_N_H	R can be alkyl, aryl. R cannot have a carbonyl carbon attached to the nitrogen of the amine.
	Amine_Sec_Alk_N_H Amine_Secondary_Alkyl_N_H	R can be alkyl. R cannot have a carbonyl carbon attached to the nitrogen of the amine. R' can be alkyl. R' cannot have a carbonyl carbon attached to the nitrogen of the amine.
	Amine_Sec_Ar_N_H Amine_Secondary_Aryl_N_H	Ar can be aryl. R can be alkyl, aryl. R cannot have a carbonyl carbon attached to the nitrogen of the amine.
	Amine_Sec_Gen_N_H Amine_Secondary_General_N_H	R can be alkyl, aryl. R cannot have a carbonyl carbon attached to the nitrogen of the amine. R' can be alkyl, aryl. R' cannot have a carbonyl carbon attached to the nitrogen of the amine.
	A_Acid_C_C Amino_Acid_C_C	R can be H, alkyl, aryl.
	Ar_or_Vinyl_Br_C_Br Aryl_or_Vinyl_Bromide_C_Br	Ar is an aryl group; Vi is a vinyl (C=C) group. The aryl or vinyl group must be directly attached to the Br.
	Ar_or_Vinyl_I_C_I Aryl_or_Vinyl_Iodide_C_I	Ar is an aryl group; Vi is a vinyl (C=C) group. The aryl or vinyl group must be directly attached to the I.
	Ar_or_Vinyl_SH_S_H Aryl_or_Vinyl_Thiol_S_H	Ar is an aryl group; Vi is a vinyl (C=C) group. The aryl or vinyl group must be directly attached to the S.

Table 4.1. Predefined functional groups with definitions

Structure	Short name Long name	Definition of R, R', R'', Alk, Ar, Vi, A
	alpha_Br_CO_C_Br alphaBromocarbonyl_C_Br	R can be alkyl, aryl. R' can be H, alkyl, aryl.
	Carbam_Cl_C_Cl Carbamoyl_Chloride_C_Cl	R can be anything with a carbon attached to nitrogen. R' can be anything with a carbon attached to nitrogen.
	alphaCarbonyl_C_H alphaCarbonyl_C_H	R can be H, alkyl, aryl, alkoxy with oxygen attached to carbonyl. R' can be H, alkyl, aryl, carbonyl with carbonyl carbon attached to CH, cyano. R'' can be H, alkyl, aryl, carbonyl with carbonyl carbon attached to CH, cyano.
	C_Acid_C_C Carboxylic_Acid_C_C	R can be anything with carbon attached to carbonyl.
	C_Acid_C_O Carboxylic_Acid_C_O	See above
	C_Acid_O_H Carboxylic_Acid_O_H	See above
	C_Ester_C_O Carboxylic_Acid_Ester_C_O	R can be anything with carbon attached to carbonyl. R' can be anything with carbon attached to the oxygen except for a carbonyl carbon.
	Cl_Formate_C_Cl Chloroformate_C_Cl	R can be anything with carbon attached to oxygen.

Table 4.1. Predefined functional groups with definitions











Structure	Short name Long name	Definition of R, R', R'', Alk, Ar, Vi, A
	Hydrazine_C_N Hydrazine_C_N	R can be H or anything with a carbon attached to the nitrogen of the hydrazine.
	Hydrazine_N_H Hydrazine_N_H	See above
	Hydrazine_N_N Hydrazine_N_N	See above
	Isocyanate_C_N Isocyanate_C_N	R can be anything with a carbon attached to the nitrogen of the NCO.
	Sulfonamide_N_H Sulfonamide_N_H	R can be anything with a carbon attached to the sulfur of the SO ₂ .
	Sulf_Cl_S_Cl Sulfonyl_Chloride_S_Cl	R can be anything with a carbon attached to the sulfur of the SO ₂ .
	Thiol_S_H Thiol_S_H	R can be H, alkyl, aryl, vinyl. R cannot have the carbon of a carbonyl attached to the sulfur.
	W_amide_C_N Weinreb_Amide_C_N	R can be alkyl, aryl. R cannot have a carbonyl carbon attached to the amide carbonyl.

Table 4.1. Predefined functional groups with definitions

Structure	Short name Long name	Definition of R, R', R'', Alk, Ar, Vi, A
	Boronic_C_B Boronic_Acid/Ester_C_B	R can be alkyl, aryl.
	Grignard_C_Mg Grignard_Reagent_C_Mg	R can be alkyl, aryl.

4.2.3 Setting Options for Structure Variation and Cleanup

In addition to defining the atoms of the reagents that will become the side chains, the reagent preparation process performs a 2D-to-3D conversion, generates all reasonable ionization and tautomeric states, and expands the stereochemistry. If the structures you have are already 3D structures in the appropriate ionization state, tautomeric and stereochemical form, you might not need to generate variations. Many collections of reagents are provided in 2D form, with implicit hydrogen atoms. To generate 3D, all-atom structures for docking, some variation to obtain structures that represent the actual state under physiological conditions may be needed.

The lower portion of the Reagent Preparation panel allows you some control over the structure variation and cleanup process, by setting options for the following:

- **Ionization state**—You can request that the structure be protonated or deprotonated to represent the actual forms in water in the given pH range, by selecting **Ionize**, using **pH** and entering the pH value and range in the accompanying text boxes. If this option is not selected, the ionization state is left as is.
- **Tautomerization state**—You can request the generation of tautomers of the input structures by selecting **Generate tautomers**. Ionized structures might have a different tautomeric form from the un-ionized structures.
- **Use of Epik for ionization and tautomerization**—Epik is a more sophisticated tool for generating ionization and tautomerization states than the default tools provided with Lig-Prep, the **ionizer** and the **tautomerizer**. It is separately licensed, so you must have an Epik license to use this option. In addition to generating ionization and tautomeric states for a given pH, it can generate states that are suitable for binding to metals. This is useful if the ligand binds to a metalloprotein.

- **Stereochemistry**—Stereoisomers of the input structures are generated in the absence of chirality information, either explicitly specified or deducible from the 3D structure. If the input structures are 2D and have chiral centers but do not have chirality information, you should ensure that the number specified in the **Generate stereoisomers** text box covers the possible combinations.
- **Ring conformations**—If the input structures contain rings that can exist in more than one low-energy conformation, enter the number of likely low-energy conformations in the **Generate low-energy ring conformations** text box. Ring conformations are not searched during docking, so the conformations must be set up during the reagent preparation process.

If you do not need to generate structural variations, you should deselect **Ionize**, using **pH** and **Generate tautomers**, and enter 1 in the **Generate stereoisomers** and **Generate low-energy ring conformations** text boxes.

4.2.4 Running the Reagent Preparation Job

When you have completed the tasks above, enter a name in the **Job name** text box and click **Run**. You can also click the **Settings** button and choose **Job Settings**, to select job options and start the job. See [Section 1.5 on page 6](#) for details.

The output structure file contains the prepared reagents. There might be several output structures per input structure if the molecules can exist in several low-energy forms. Consequently, there can be multiple structures for a given combination of reagents on a core (a “compound”) and the final number of structures in the library can be greater than the final number of compounds.

The following output files are required for later use:

<code>jobname.bld</code>	Reagent structures in all modifications labeled for building
<code>jobname.log</code>	Overall log file for the job
<code>jobname.sqlite</code>	Sqlite database built from the .bld file
<code>jobname_vpost-stats.txt</code>	Text file giving summary statistics of molecules in the .bld file

4.2.5 Task Summary

To set up and run a reagent preparation job:

1. Choose the source of reagent structures
2. Select the functional group
3. Select 2D-to-3D conversion and structure variation (LigPrep) options
4. Click Run.

4.3 Preparing the Core-Containing Molecules

The structures that you select for the core-containing molecules must be all-atom, minimized, 3D structures. If the structure you have does not meet these criteria, you can prepare it using LigPrep. For more information, see the [LigPrep User Manual](#).

Note that only the ionization state and tautomer that you select is used for the core-containing molecule when you define the core in the Define Combinations step of the Combinatorial Screening panel or in the Combinatorial Library Enumeration panel. If you want libraries with multiple states of the core, you must generate the states of the core-containing molecule and select them one-at-a-time to generate a library based on each state.

Combinatorial Screening: The Process

The goal of combinatorial screening is to create a small combinatorial library of structures that are likely to have high activity towards to a selected target. The library is screened, first by docking to the target receptor and eliminating structures that do not dock well, then by analyzing the docked structures to define a small reagent set that is likely to have a large number of actives among the compounds that are generated.

Once you have decided on the chemistry, obtained files containing the reagents, the receptor and the core-containing molecule, and prepared the structures in these files, you can proceed with the combinatorial screening process. This process is managed from the Combinatorial Screening panel. To open the Combinatorial Screening panel, choose Applications → CombiGlide → Combinatorial Screening.

5.1 The Combinatorial Screening Panel

The Combinatorial Screening panel is designed like a wizard, with five steps. To perform a combinatorial screening, you proceed through each step in turn. You do not have to complete all steps in a single session. The design allows you to exit at an intermediate stage and pick up the process again later at any step for which you have the required data.

Each set of data defined by proceeding through the steps in this panel is called a *run*. The results for each run are stored as a separate entity. You can open existing runs, and create and save runs from the File menu. Runs are saved in the Maestro project, so if you want to keep your results, you should create a named Maestro project to store them. See [Section 9.1](#) of the *Maestro User Manual* for more information.

If you backtrack in a given run and make changes, the results for all the steps that depend on what you changed are discarded. When you make a change that affects later results, you are prompted to save the previous results and create a new run.

The step features occupy the center of the panel, and consist of a title with a brief description of the step at the top, a set of controls and tables for results, and a Back and a Next button at the bottom. The tasks for each step are summarized in the next section, and the steps are described in detail in the following chapters.

In addition to the step features, the panel contains a menu bar and an job status button at the top, and a step guide at the bottom. These features are described below.

Table 5.1. Description of the File menu.

Menu Item	Description
New	Create a new run. Opens a dialog box to specify the run name. The new run becomes the current run.
Open	Open an existing run from the submenu. If there are more than 4 runs, choose More to open a dialog box and select a run.
Save As	Save the current run with a new name. Opens a dialog box to specify the new run name. The run is copied, and the renamed run becomes the current run.
Rename	Rename the current run. Opens a dialog box to specify the new run name.
Delete	Delete the current run.
Import	Import the results of a run that was initiated from the command line. Prompts for a run name, then opens a file selector in which you can browse to a directory, display the available runs, and select a run.

The File menu

The File menu allows you to work with the runs that are available in the project. The items on this menu are described in [Table 5.1](#).

The Step menu

The Step menu contains an item to display or hide the Guide, and an item for each of the steps. The current step is marked as selected; likewise if the Guide is displayed, it is marked. The items for the steps that are not available are dimmed. You can go to any available step by choosing the corresponding menu item.

The job status button

When a job has been launched and is running, the icon on the job status button at the upper right of the panel starts spinning. When the job finishes, the icon stops spinning. If incorporation of results takes some time, the icon is replaced by an exclamation point while the incorporation is taking place. To monitor the job using the Monitor panel, click the button. For more information about monitoring jobs, see [Chapter 3](#) of the *Job Control Guide*.

The Guide

The Guide displays the steps in the model as a set of buttons linked by lines. The buttons for the steps that are not available are dimmed. The current step is highlighted with a white background. You can go to any available step by clicking its button in the Guide. The Guide can be displayed or hidden from the Step menu.

5.2 The Combinatorial Screening Process

The tasks involved in each step of the combinatorial screening process are summarized below.

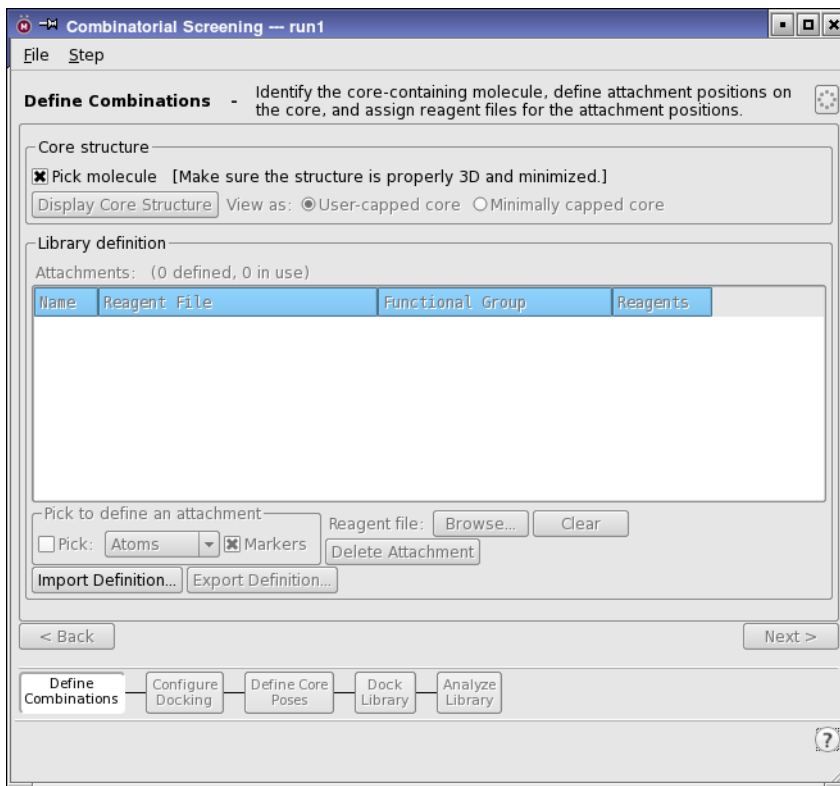


Figure 5.1. The Combinatorial Screening panel.

The Define Combinations step

In this step, you first select the core-containing molecule that you want to use for generation of the library. The library is generated by replacing selected side chains in the molecule with other side chains from the reagent files. For each side chain, you identify the bond that defines where the substitution will take place (the *attachment position*), and specify a reagent file that contains the reagents with the side chains that you want to substitute. The bond is called the *grow bond*. The combination of an attachment position and a reagent file is called an *attachment*. The core of the molecule is the part of the molecule that is left when you remove the side chains. The reagent libraries must be prepared beforehand, using the Reagent Preparation panel.

Once you have selected a molecule and the attachment positions for the reagents, and added the reagent files, you can proceed to the next step.

The Configure Docking step

The screening process uses CombiGlide docking to filter out reagents that produce molecules that do not score well. This step assumes that you have already generated the necessary Glide grid files, using the Receptor Grid Generation panel. In this step, you set options for the docking of the libraries to the receptor. As well as selecting a grid and making basic settings as for a Glide run, you can specify the full range of Glide constraints for the docked molecules.

After selecting a grid and setting any desired options, you can proceed to the next step.

The Define Core Poses step

In this step, you select the method for determining the poses of the core structure that you want to use in the docking step. These poses are generated by docking one or more molecules that contain the core. These molecules can be the core-containing molecule, the minimally capped core, or some other molecules that you choose. The *minimally capped core* is the core of the molecule with a minimal capping group at each attachment position. In most cases, the minimal capping group is defined by R, R', R''=Me or Ar=Ph for each reagent. See [Table 4.1 on page 29](#) for definitions of R, R', and R''.

You can constrain the core position to prevent it from moving to regions of the receptor that would generate unacceptable binding modes.

When you have selected a method, you can proceed to the next step.

The Dock Library step

In this step you run the docking jobs. The docking is performed in three stages. In the first stage, the core-containing molecules are docked to obtain a set of core poses, using whatever constraints on the core position you imposed.

In the second stage, each reagent from each reagent library is added in turn to the core, to produce molecules with a single side chain and the minimal capping groups at the remaining attachment positions. These molecules are docked using CombiGlide XP docking. Reagents that do not generate any good poses are eliminated. This stage can be run independently, and is known as the “single-position docking” stage.

In the third stage, the poses with a single side chain are combined intelligently to produce molecules that are likely to dock well. These molecules have side chains from the reagent libraries at all attachment points. This stage is the most time-consuming part. Glide constraints are only applied in this stage. This stage is known as the “combinatorial docking” stage.

After docking, you can export the results of either the single-position docking or the combinatorial docking for further examination. From this step you can also enumerate and dock the entire library without performing the initial screening.

Once you have docked the combinatorial library, you can proceed to the next step.

The Analyze Library step

In this step, the number of reagents used at each position is reduced to a small set by application of filtering and selection strategies. Filtering is based on structural properties such as molecular weight and counts of structural features, and ADME properties. These properties are generated by QikProp, which is run automatically. The selection strategies identify the “best” reagents at each position, based on the GlideScore of each of the docked structures. When the small reagent set is selected, the library defined by this set can be enumerated.

Combinatorial Screening: Defining the Core and Its Attachments

The combinatorial screening process begins with the definition of the core and its attachments. The core-containing structure must be a 3D, all-atom, minimized structure. Likewise, the reagents used for the attachments must be 3D, all-atom, minimized structures, and also include information that defines the side chain to be attached to the core. Preparation of the core-containing molecule and the reagents is described in [Chapter 4](#). The task of selecting the core-containing molecule, defining the attachment points, and associating a set of reagents with each attachment point is carried out in the Define Combinations step of the Combinatorial Screening panel.

6.1 Selecting the Core-Containing Molecule

The first task in this step is to select the core-containing molecule. When you first enter the Define Combinations step, Pick molecule is selected in the Core structure section, and most other controls are unavailable (see [Figure 6.1](#)).

To select the core-containing molecule:

1. Ensure that the desired molecule is displayed in the Workspace.

You can import the structure if necessary, or open the Project Table panel to display the structure, while the Combinatorial Screening panel is open. If you import the structure, ensure that Replace Workspace is selected.

2. Click on an atom in the core-containing molecule.

Any structures other than the core-containing molecule are undisplayed. The other controls in the step become available.

The molecule that you picked is copied into the CombiGlide run. Any changes you make to the structure of the molecule are not saved in the run. If you make changes to the structure, you must pick the molecule again to make it the core-containing molecule. If you have already picked a molecule, picking a new molecule deletes all information about attachments.

After you pick the core-containing molecule, Pick molecule is deselected, and most other controls become available. The Pick option is selected in the Pick to define an attachment section. The core-containing molecule is now defined, and you can define the positions of the attachments, which also defines the actual core.

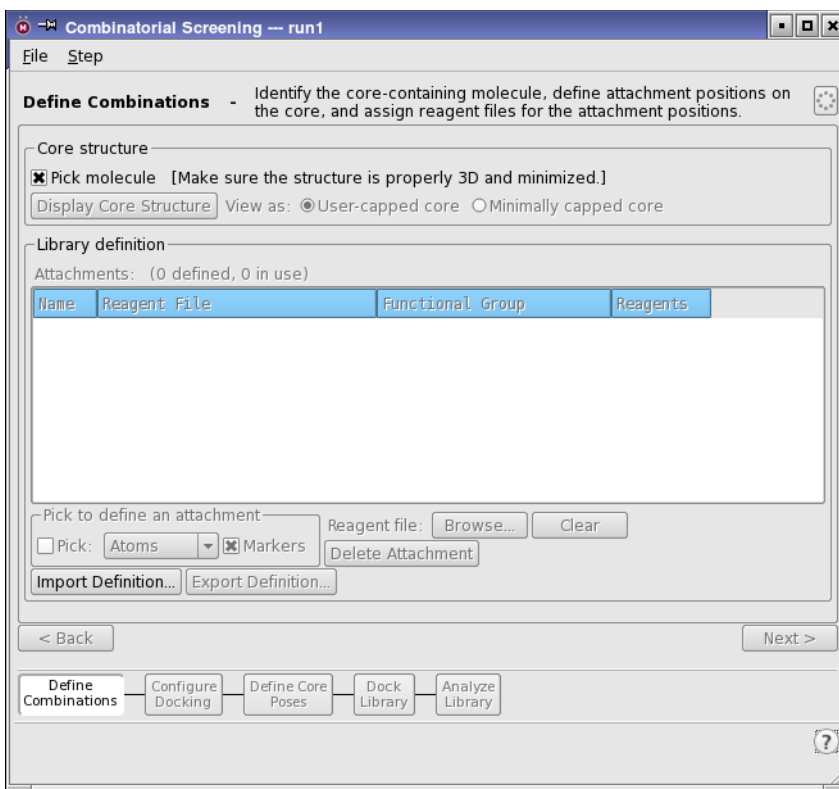


Figure 6.1. Initial view of the Define Combinations step.

6.2 Defining Attachments

Once you have a core-containing structure, you can define the attachments. An *attachment* is defined as the combination of an attachment position and a file that contains the reagents that will be attached at this position. The attachment position is defined by the atom in the core to which a side chain is attached and the atom in the side chain that is bonded to it. To make the definition, you must select both the atom that is kept in the core, and the atom that is bonded to it in the side chain. The order in which you pick these atoms determines which is the core and which is the side chain. The bond so defined is called the *grow bond*: it is essentially a vector that points from the core to the side chain. The direction is the direction in which the side chain is “grown” from the core.

Having defined the attachment position and the side chain, the next task is to select a set of reagents for this attachment position from which the side chains are taken. The reagents are contained in files that you prepared using the Reagent Preparation panel.

To define an attachment:

1. Pick two atoms in the core-containing molecule that define the attachment position, the atom that is kept first, then the atom that is deleted.

Alternatively, choose Bond from the Pick option menu in the Pick to define an attachment section, and click on the appropriate bond. When you click, make sure you click on the end of the bond closest to the side chain, to ensure that the grow bond points in the right direction. This is especially important for the first attachment. For subsequent attachments, a warning is posted if the grow bond points in the wrong direction.

After picking the second atom (or the bond), a gold arrow is displayed over the bond, and the Select Reagent File dialog box opens.

2. Select the appropriate reagent file, and click OK.

The reagent files have a .bld extension: they are Maestro-format files that have special information for the Maestro builder, which does the library enumeration. The dialog box closes, and the information in the row of the Attachments table is completed. The size of the library with the current set of attachments is given in the lower portion of the panel.

The arrowhead is replaced with a gold sphere centered on the side chain atom.

If you choose the wrong location or the wrong direction for an attachment, you can delete it by selecting it in the table and clicking Delete Attachment. When the error is recognized by CombiGlide, such as if you pick a bond in a ring, a warning is posted and the attachment is cleared. If you simply attached the wrong file, you can select the attachment in the table, and click the Browse button below the table to select the correct file, or click Clear to remove the attachment file.

For each attachment a minimal capping group is automatically defined. This minimal capping group is usually a small instance of the side chain in which R, R', R''=Me or Ar=Ph for the chosen reagent type. Adding the minimal capping group to the core defines the *minimally capped core*. You can view the minimally capped core by selecting Minimally capped core from the View as options in the Core structure section. To redisplay the original core-containing structure, select User-capped core.

Information for the attachments that have been defined is listed in the Attachments table. The table columns are described in Table 6.1. You can select multiple rows in the table with the usual shift-click and control-click actions. The markers for the selected attachments are colored turquoise in the Workspace.

Table 6.1. Description of Attachments table columns

Column	Description
Name	Label for the position of the attachment. Default is 1, 2, 3, etc. Editable.
Reagent File	Name of the file containing the reagents, minus the extension. The tooltip for this table cell displays the full path to the file.
Functional Group	Long name of the functional group that describes the reagents.
Reagents	Number of reagents in the file.

Right-clicking in a table row displays a menu with four items: Select Reagent File, Clear Reagent File, Reagent File Statistics, and Delete Attachment. The actions are applied to all selected rows. The Reagent File Statistics item displays information on the selected reagent files from the reagent preparation process, including the number of structures that were generated at each step and molecular weight data statistics. The other items perform the same actions as the Browse, Clear, and Delete Attachment buttons.

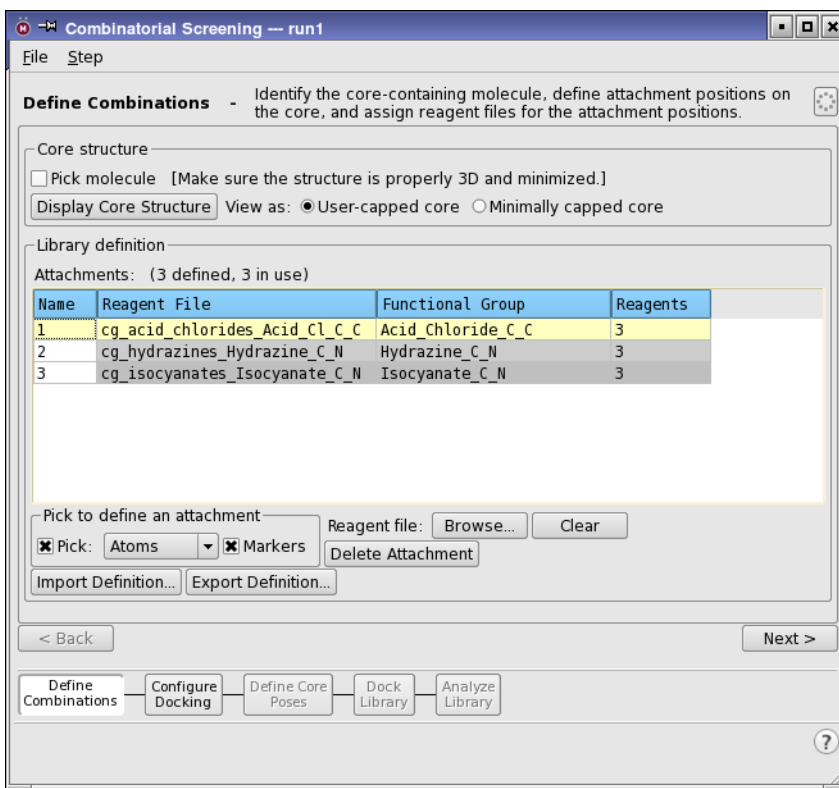


Figure 6.2. The Define Combinations step after defining attachments.

Note: If you define only one attachment point, the only workflow available to you is fully enumerated docking. The postdocking library analysis depends on the existence of more than one attachment point.

6.3 Saving and Loading Library Definitions

The combination of a core structure and its attachments defines a combinatorial library. You can save this definition for future use, by clicking **Export Definition**. A file browser opens, in which you can navigate to the desired location and specify a name. The library definition is saved as a gzipped tar file, and the extension `-comdef.tar.gz` is added to the name automatically. The library definition includes the core structure, which is written to the file `core-comdef.mae`, and the files for each of the attachments. These files are all added to the tar file.

If you want to use a previously-defined library definition, including the core structure and the attachments, click **Import Definition**, and navigate to the desired gzipped tar file. The default filter for the file selector is `-comdef.tar.gz`, so you should only see valid library definitions.

Combinatorial Screening: Setting Up for Docking

The center of the combinatorial screening process is the docking of the library members using CombiGlide XP docking. After you have defined the library, the next task is to set up the parameters and options for docking.

The docking of the library is performed in three stages:

- **Core docking.** In this stage, the core-containing molecules are docked to determine the best core poses. By default, the user-capped core is used. These poses are used as the initial poses for the substituted structures that constitute the library members. Two jobs are run in this stage, called `core` and `mcore` docking.
- **Single-position docking.** In this stage, the side chains at each position are added to the core, one at a time, to generate a set of structures with side-chain substitution at a single position. The other attachment positions are capped with either the minimal capping group or the user-defined capping group for that position. These singly-substituted structures are then docked, and passed through a selection process that eliminates reagents that are likely to result in poor binding.
- **Combinatorial docking.** The final stage is to dock a selected number of fully substituted structures that were predicted to have the highest probability of binding to the receptor by the selection algorithm.

Setting up the docking calculations is done in two steps: a configuration step, in which Glide parameters and constraints are set, and a core pose definition step, in which the structures to use for the core docking are selected and any constraints on the core position are set. These tasks are performed in the Configure Docking and Define Core Poses steps of the Combinatorial Screening panel.

7.1 Making Glide Settings

The Glide settings are made in the Settings tab of the Configure Docking step. This tab is a combination of parts of the Settings tab and the Ligand tab from the Glide Ligand Docking panel. Because CombiGlide uses Glide XP docking and the ligands are predefined, only the relevant options from these tabs are available. There are three sections in this tab: Receptor grid, Docking, and Van der Waals radii scaling.

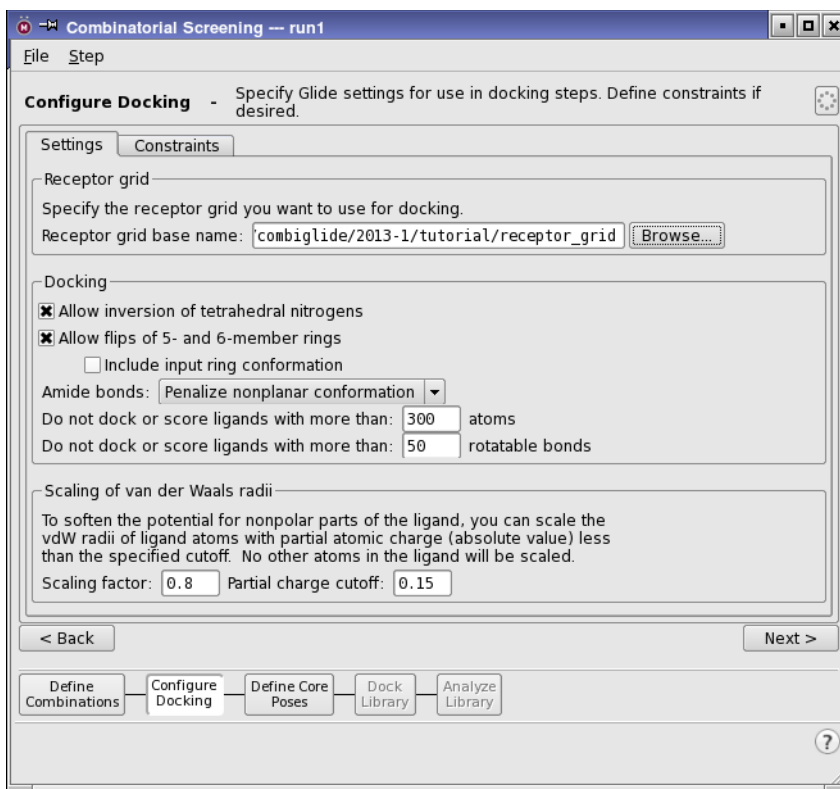


Figure 7.1. The Settings tab of the Configure Docking step.

The first task is to select the receptor grid. You can either enter the path to the receptor grid file in the Receptor grid base name text box, or click Browse and navigate to the receptor grid file.

If you have not generated a grid, you must do so before proceeding, in the Receptor Grid Generation panel. You can open this panel from the CombiGlide submenu of the Applications menu in the main window. When you have generated a grid, you can reopen the Combinatorial Screening panel and return to this step.

Next, you can set options and parameters as listed below. The ring flip and amide bond rotation options are applied in all docking stages. The limits on the maximum number of atoms and rotatable bonds are only applied in the combinatorial docking stage. Structures exceeding these limits are not built and thus are not docked.

For more information on these options, see [Chapter 5](#) of the *Glide User Manual*.

Allow inversion of tetrahedral nitrogens

Select this option to allow inversion at tetrahedral nitrogens. This option is selected by default.

Allow flips of 5- and 6-member rings

Select this option to allow the conformation of nonaromatic 5- and 6-membered side-chain rings to be varied during docking. This option is selected by default. If you select this option, you can also select Include input ring conformation to ensure that the input ring is included; otherwise it is only included if it is sufficiently low in energy.

Amide bonds

Choose an option from this option menu for the treatment of amide C–N bonds. The options are Penalize non-planar conformation, which penalizes amide bonds that are not cis or trans when scoring the docked molecules, Vary conformation, which allows nonplanar amide bonds, Retain original conformation, which freezes amide bonds in their input conformation throughout docking, and Allow trans conformation only, which enforces trans amide bonds to within a small angle range. The default is to penalize nonplanar amide bonds.

Do not dock or score ligands with more than n atoms

Enter a number to restrict the size of the molecules that are docked. The default is 300. This restriction is only applied in the combinatorial docking stage.

Do not dock or score ligands with more than n rotatable bonds

Enter a number to restrict the number of rotatable bonds to sample during docking. The default is 50. This restriction is only applied in the combinatorial docking stage.

Van der Waals radii scaling

Enter a scaling factor and a partial charge in the text boxes to scale the van der Waals radii of atoms with small partial charges. The scaling simulates the flexibility of nonpolar groups.

7.2 Setting Glide Constraints

Constraints can also be used with CombiGlide and are specified in the same manner as with Glide. With CombiGlide, constraints are only turned on during the combinatorial docking stage, i.e. docking of the fully substituted structures. This prevents the docking algorithm from forcing the side chains at different positions on the core into the same binding pocket during the single-position docking. In the combinatorial docking stage, constraints can be used to filter out reagents whose side chains cannot satisfy the relevant constraints.

The Constraints tab of the Configure Docking step tab contains controls for applying Glide constraints to the docking of the combinatorial library. The contents of this tab are identical to those of the Constraints tab in the Glide Ligand Docking panel. For information on setting constraints, see [Section 5.6](#) of the *Glide User Manual*.

7.3 Defining the Core Poses

The final setup task is to determine how the core poses will be generated in the first stage of the docking step. The core poses can be generated with any molecule that includes the core, not just the structure you selected to define the core in the Define Combinations step. This task is carried out in the Define Core Poses step, and involves selecting a docking mode and specifying the structures that will be docked.

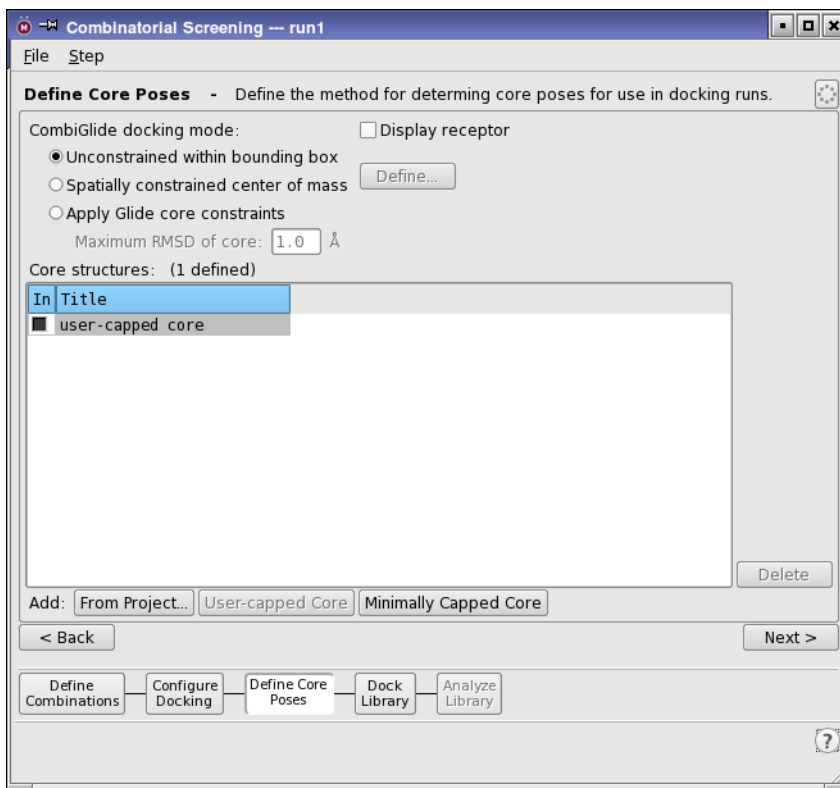


Figure 7.2. The Define Core Poses step.

7.3.1 Selecting a Docking Mode

The first stage in the docking step is to dock the core-containing molecules. This stage does not apply any Glide constraints that you have set. If the molecules used to define the core poses are smaller than the original core-containing molecule, they might bind to locations in the active site that are not accessible to the original. Many of these poses could be undesirable, for example if the core migrated to a location in which the fully substituted molecules could not dock. CombiGlide therefore supplies three docking modes, described below, two of which allow some kinds of constraints to be applied.

Unconstrained within bounding box

This is the default option and permits the molecule the widest range of movement within the grid bounding box during docking. If the molecules you select for the core poses are likely to dock in reasonable poses, select this option.

Spatially constrained center of mass

With this option, the core center of mass is constrained during docking to a sphere of a given radius. Constraining the center of mass allows the core-containing structures to rotate but prevents them from moving to another possible binding site.

The position of the sphere is defined by picking atoms in the Workspace, and you can choose atoms in the ligand or the receptor, or both. If you want to display the receptor in order to pick atoms, click Display Receptor. You can also display any of the core-containing structures by clicking the appropriate In column of the Core structures table. However, you should make sure that these structures are aligned to the receptor, otherwise picking atoms from them will be of no value. When you have the desired structure in the Workspace click Define. In the Define Core Constraint dialog box, use the picking tools to pick the atoms whose centroid defines the sphere center.

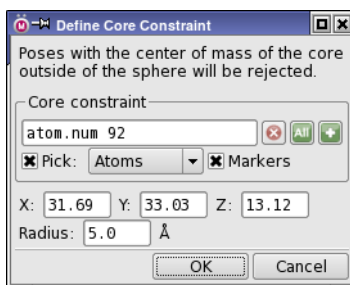


Figure 7.3. The Define Core Constraint dialog box.

The radius of the sphere is 5.0 Å by default, and can be set in the Define Core Constraint dialog box. The sphere radius must not extend beyond the enclosing box of the grid. You can also set the coordinates of the center of mass manually.

Apply Glide core constraints

This option allows you to define the core poses by the positions of the core-containing structures, and provides the greatest degree of constraint on the core poses. These structures must be already aligned to the receptor. When these structures are docked, poses for which the RMSD of the core atoms exceeds the threshold specified in the Maximum RMSD of core text box are rejected. One way of obtaining these core-containing molecules is to dock molecules that contain the core onto the same receptor using Glide, with whatever constraints you want to apply. Another is to use the crystal structure of the chosen receptor complexed with a core-containing molecule.

7.3.2 Specifying Structures for the Core Poses

The structures that will be docked to define the core poses are listed in the Core Structures table. You can add structures to the table with the Add buttons below the table, and you can delete structures from the table by selecting them and clicking Delete.

By default, only the user-capped core is listed. This is the core-containing molecule that you selected in the first step. If you delete this molecule, you can add it again by clicking User-capped Core. You can also add the minimally capped core to the table by clicking Minimally Capped Core. To add structures from the Project Table, click From Project, and select the molecules in the entry chooser that opens. The entries must satisfy the following criteria:

- They must contain only a single molecule
- The molecule must contain the core structure

No checking is done for duplicates, so you must ensure that you don't add the same structure more than once. The structures are copied into the run, so that any changes in the Project Table do not affect these structures.

If a crystal structure or docked pose is used for the core structure, the associated protein structure must be superimposed on the receptor structure from which the grids were calculated to ensure that the core poses are in the correct frame of reference. You can superimpose the protein structures in the Superposition panel, or use the Align Binding Sites panel, both of which you open from the Tools menu in the main window.

Combinatorial Screening: Docking the Library

The docking step is the heart of the combinatorial screening process, where the docking jobs are run to screen out molecules that do not dock well and thereby eliminate reagents. This step is carried out in the Dock Library step of the Combinatorial Screening panel.

First, the core-containing structures are docked. The resulting poses are clustered, and poses from each cluster are selected until the maximum number is reached. This ensures the widest diversity of core poses while limiting their number. The core poses are used as the initial positions of the substituted molecules in the subsequent docking stages.

The screening is performed in two stages. In the first stage, structures that have only a single substitution are docked. These structures are generated by adding the reagents, one at a time, to the minimally capped core or the user-capped core. The result is a set of molecules (a “single-position library”) for each attachment position with a side chain at that position and the minimal capping group at the other positions. The side chains added are taken from the expanded set of reagents, in which different ionization states, stereoisomers, conformers, or tautomers can be represented. In the pyrazole example, there are three attachment positions and thus, three single-position libraries will be generated and docked. This portion of the virtual library evaluation is an additive process: the number of dockings required for this step is the sum of the number of side chains for each position.

Once the single-position docking results are returned, a selection algorithm determines which fully substituted cores have the highest probability of binding to the receptor. Structures from the single-position docking stage that do not dock well are eliminated. Reagents for which no structures docked well are removed from the reagent list for the appropriate attachment position. The selection algorithm takes into consideration the XP GlideScores of the docked poses from the single-position docking. It also checks that the side chains it has identified for a particular fully substituted core are not predicted to compete for the same part of the receptor binding site.

Finally, the structures that are identified by the selection algorithm are generated and docked. This set of structures is not a combinatorial library, in which each side chain occurs at each position.

As an alternative, the single-position docking and the selection procedure can be bypassed, and the entire library can be enumerated and docked.

8.1 Running the Docking Jobs

The docking jobs are set up and run from the Dock dialog box, which opens when you click Dock. This dialog box offers choices for the docking task and options for each of these tasks.

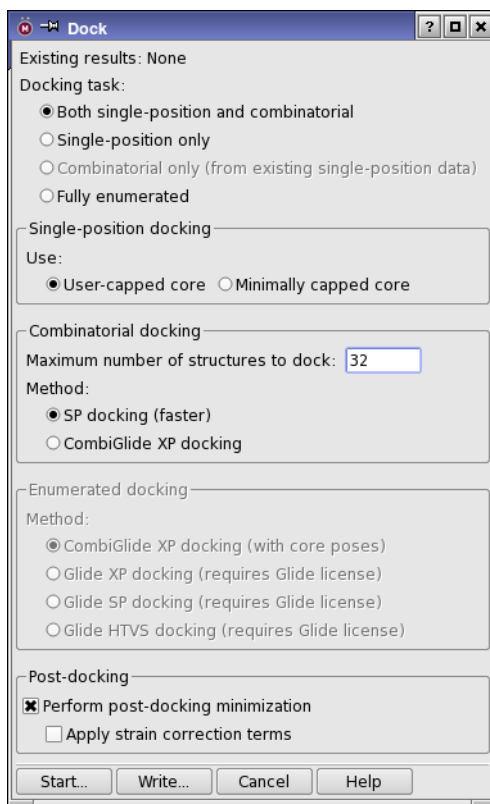


Figure 8.1. The Dock dialog box.

If you want to continue on in the workflow to analyze the docking results and select the optimal library, you need to run both the single-position docking and the combinatorial docking stages. These two stages can be run in the same job or separately. If you want to examine the results of the single-position docking before continuing with the combinatorial docking, select Single Position only. When you want to proceed to the second stage, open the Dock dialog box again, and select Combinatorial only. The results of the first stage are then used to run the combinatorial docking. To run both stages in the same job, select Both single-position and combinatorial. The core docking is run automatically before the single-position docking. The status of any existing results is displayed at the top of the dialog box, to help you choose the appropriate task.

Once you have chosen the task, you can choose options for the two docking stages. For single-position docking, you can use either the user-capped core or the minimally capped core as the molecule on which the single substitutions are made. The benefit of the user-capped core is that the side chains on the unsubstituted positions are usually larger and may prevent docking in unreasonable poses. Using the minimally capped core allows greater flexibility in docking poses, and might identify poses that would otherwise be missed.

For the combinatorial docking, you can set a limit on the number of fully substituted structures to dock in the Maximum number of structures to dock text box. The default is 20 times the sum of the number of side chains at each position. If you plan to use Glide constraints, you should significantly increase the limit to ensure sufficient output structures. You can also choose whether to use SP docking or CombiGlide XP docking.

If you want to dock the entire combinatorial library, select Fully enumerated for the docking task. You can then select the Glide docking mode. In addition to using CombiGlide XP docking, you can select any of the Glide docking modes: XP, SP, or HTVS. You must have the appropriate Glide license for Glide docking. If you have a single attachment point, this is the only option that is available.

For any of the tasks, you can also choose to perform a post-docking minimization and adjust the scoring for ligand strain, in the Post-docking section. Post-docking minimization usually results in better poses, because it adjusts bond lengths and angles as well as dihedrals. The ligands are rescored after the post-docking minimization, with the strain correction if it included. For more information, see [Section 5.8.2](#) of the *Glide User Manual*.

When you have chosen all the options, click Start, set any job options in the Start dialog box, such as selection of a host and distribution of the job over multiple processors. Click Start to submit the job.

When the combinatorial docking stage finishes, QikProp is run for the completed structures, and the QikProp properties are added to the run. These properties can be used in the Analyze Library step.

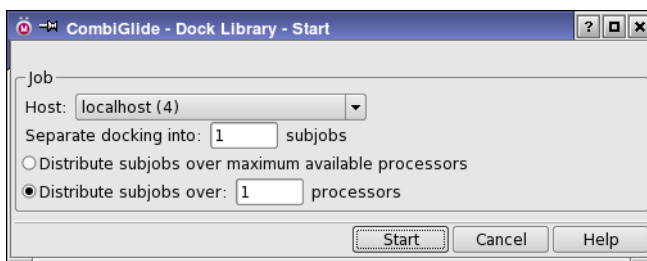


Figure 8.2. The Start dialog box.

To cancel the current docking job, click **Cancel**, then click **OK** in the confirmation dialog box. If you are running the combinatorial docking job, and the single-position docking job has already finished, the results of the single-position docking job are retained.

8.2 Docking Results

Statistics on the reagents that survive the various docking stages are added to the **Reagent retention** table when the docking jobs are finished. The table is noneditable, and its columns are described in [Table 8.1](#). The number of structures actually generated is reported in the **Combinatorial results** section. This number can be less than the product of the numbers in the **Comb** column of the **Reagent retention** table, because of the limits imposed and because not all structures generated from the single-position docking necessarily dock well.

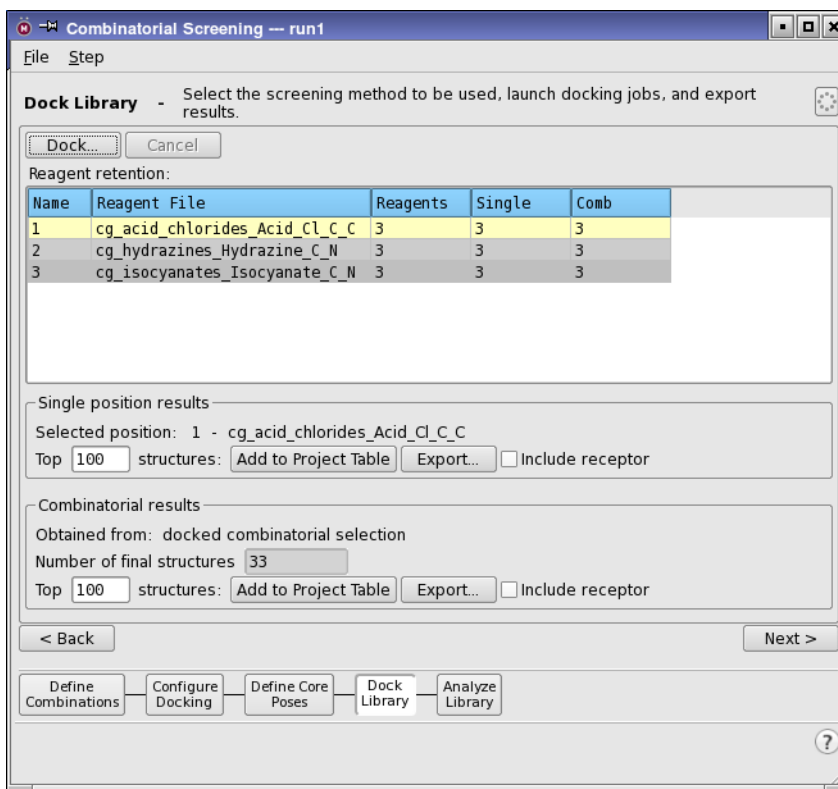


Figure 8.3. The **Dock Library** step after docking.

Table 8.1. Description of Reagent retention table.

Column	Description
Name	Attachment position label
Reagent File	Name (minus extension) of the file that contains the reagents for the specified position.
Reagents	Number of reagents in the original reagent file.
Single	Number of reagents from this file that survived the single-position docking.
Comb	Number of reagents from this file that survived the combinatorial docking. This number includes all reagents that contributed to any structure that docked well.

Once the results are available, you can add them to the Project Table or export them to a file. The structures from both stages are sorted by their GlideScores. Because there may be variations of each reagent (such as different ionization states, tautomers, conformers), it is the top-scoring variation for each reagent that is exported. For the single-position docking results, multiple poses are included in the set that is exported: thus the top-scoring variation for each pose for each reagent is exported. The set of singly substituted or the set of fully substituted structures is added to the Project Table as an entry group, with their QikProp properties. If the receptor is included, it is the first entry in the entry group or the file.

To export a set of structures from single-position docking:

1. Select the row in the Reagent retention table for the desired reagent type.
2. Enter the number of reagents in the Top *n* structures text box in the Single position results section.
3. (Optional) If you want to include the receptor, select Include receptor.
4. Click Add to Project Table or Export.
5. If you are exporting to a file, navigate to the directory, enter the file name, and click Export.

To export a set of structures from combinatorial docking:

1. Enter the number of structures in the Top *n* structures text box of the Combinatorial results section.
2. (Optional) If you want to include the receptor, select Include receptor.
3. Click Add to Project Table or Export.
4. If you are exporting to a file, navigate to the directory, enter the file name, and click Export.

Combinatorial Screening: Focusing the Library

The final step in the combinatorial screening process is to focus the library to a small number of compounds generated from a few reagents at each position. This focused library should contain a large proportion of compounds that are likely to bind well. The focusing is carried out by applying various selection strategies, by filtering on molecular and predicted ADME properties, by manual selection, and by setting limits on the number of reagents at each position and on the overall library size. The purpose of this step is to select sets of optimal reagents for each attachment position that are small enough and chosen well enough for practical synthesis and screening of the resulting real combinatorial library.

The selection and filtering can be repeated as many times as you like, and the results of each selection and filtering job can be stored and compared. You can combine selection, filtering and manual inclusion or exclusion of reagents. When you are satisfied with a selection, you can run a job to create and save the structures in the library.

The focusing of the library is driven from the Analyze Library step of the Combinatorial Screening panel. The selection and filtering tasks are set up and run from the Filter and Select dialog box, which you open by clicking Filter and Select. This dialog box is nonmodal, so you can interact with the Workspace and the Combinatorial Screening panel while it is open.

9.1 Selection Strategies

CombiGlide has three automated selection strategies, which can be supplemented with manual selection. The automated strategies use the GlideScore to select the reagents that produce the greatest number of structures that have a good GlideScore.

The selection strategies work by filling lists of reagents at each position, until the limits on the minimum and maximum number of reagents and the overall library size are met. The molecules from the docking step are ordered according to their GlideScore or some derived quantity. Each molecule is considered in turn, and the reagents that were used to build that molecule are added to the corresponding lists.

The strategies are selected in the Selection Strategy tab of the Filter and Select dialog box.

The first strategy selects reagents based on the single-position docking. In this strategy, the lists for each position can be filled independently, from the docking results for each position. This is the only strategy available if you do not have combinatorial docking results. To apply this strategy, select Single position under Selection strategy, and click Apply or OK.




Figure 9.1. The Selection Strategy tab of the Filter and Select dialog box.

The second strategy selects reagents based on the GlideScores from the combinatorial docking. Starting from the structure with the best GlideScore, the reagent for each position is added to the list for that position. The reagents for the next best structure are then added to the lists, unless they have already been added, and so on, until the conditions on the minimum and maximum reagents at each position and on the maximum library size are met. When a list at a particular position is filled, there are two options for adding structures to the remaining partially filled lists. The first option is to add reagents from only those molecules that have reagents at the filled position that are in the filled list. The second option is to add reagents from any molecule, regardless of whether the reagent at the filled position is in the filled list. The first option ensures that all molecules used to determine the library are in the library. The second option ensures that the reagents came from the structures with the best GlideScores. To apply this strategy, select Combinatorial under Selection strategy, then select GlideScore. If you want to ensure that the library contains the molecules used to determine it, select Guarantee inclusion in library. This option is only available if you have minimum and maximum numbers in the Restrictions table (see below). Click Apply or OK to start the selection job.

The third strategy uses the GlideScore from the combinatorial docking as a free energy, and calculates the normalized Boltzmann population for each docked molecule. The populations are summed for each reagent, and the reagent populations are converted to a relative GlideScore using a free-energy relation. Reagents are added to the lists in order of relative GlideScore, from lowest to highest. This strategy samples the most frequently occurring reagents,

weighted according to GlideScore: reagents coming from molecules with good GlideScores are weighted more heavily than those with poor GlideScores. For each position, the list of reagents is filled, starting from the reagent with the lowest relative GlideScore, until the conditions on the minimum and maximum reagents at each position and on the maximum library size are met. To apply this strategy, select **Combinatorial** under **Selection strategy**, then select **Boltzmann population**. Click **Apply** or **OK** to start the selection job.

In addition to applying a selection strategy, you can set the minimum and maximum number of reagents at each position, and set an overall library size. The maximum library size can be entered in the **Maximum library size** text box. The minimum and maximum number of reagents at each position can be set by editing the table cells in the **Restrictions** table, which is described in [Table 9.1](#).

Table 9.1. Description of the Restrictions table.

Column	Description
Name	Attachment position label
Reagent File	Name of reagents file for this attachment position, without the extension.
Min	Minimum number of reagents to include at this position. Editable.
Max	Maximum number of reagents to include at this position. Editable.

9.2 Manual Selection of Reagents

You can manually select reagents that must be included in the library or that must be excluded from the library in the **Include/Exclude** tab of the **Filter and Select** dialog box. This allows you to enforce the selection of reagents that did not score well, or exclude reagents that did score well, based on your knowledge of the chemistry, or to exclude reagents based on their availability.

To select reagents for inclusion or exclusion:

1. Select an attachment position from the **Position** option menu.
All the reagents in the reagent file for that position are displayed in the **Reagents** list.
2. (Optional) Filter the reagents by entering a string in the **Filter** text box.

You can use the "*" character to represent an arbitrary string. All reagents whose titles contain the specified string are listed in the **Reagents** list when you press **ENTER**.

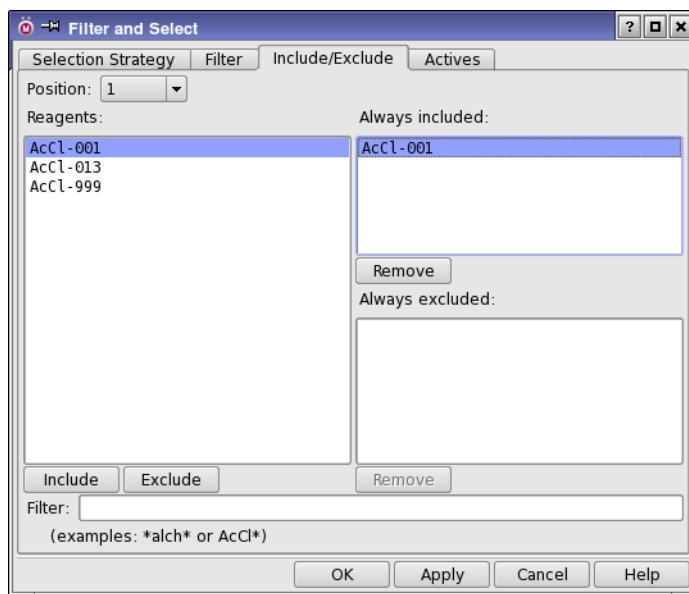


Figure 9.2. The Include/Exclude tab of the Filter and Select dialog box.

3. Select reagents from the Reagents list.
4. Click Include to add the selected reagents to the Always Included list, or click Exclude to add the selected reagents to the Always Excluded list.

If you add a reagent to one list that is already in another list, it is removed from the other list.

To remove reagents from a list:

1. Select the reagents in the list.
2. Click the Remove button immediately below the list.

You can include or exclude reagents directly from the Analyze Library step, by right-clicking a reagent in the Reagents at position *N* table. When you do so, the Filter and Select dialog box opens so that you can re-run the library selection job with the newly included or excluded reagent.

9.3 Filtering

In addition to selecting reagents with one of the selection strategies, you can filter the combinatorial results based on various predicted ADME properties of the molecules. These properties are generated by QikProp. Several filters are provided, and you can customize these filters by including or excluding any of the listed properties and by setting the limits on the acceptable range of the properties. The property ranges in the filters represent the undesirable property ranges: molecules whose properties fall inside these ranges will be filtered out. The filters are set up in the Filter tab of the Filter and Select dialog box.

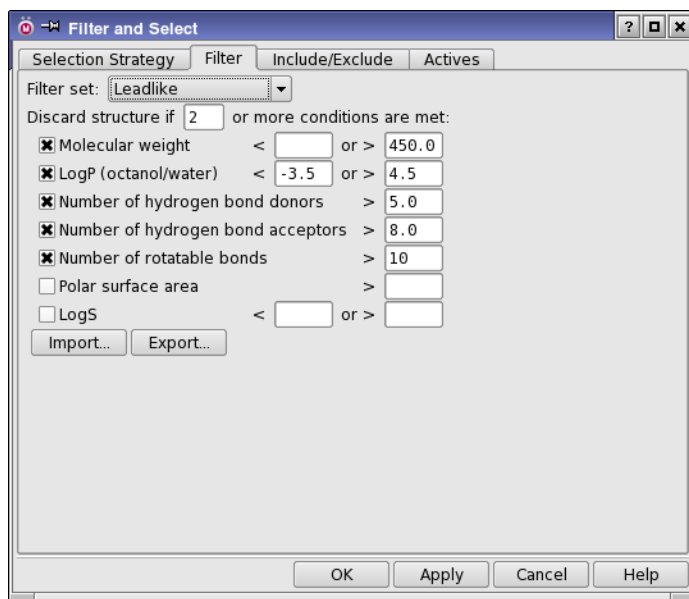


Figure 9.3. The Filter tab of the Filter and Select dialog box.

If you apply both a filter and a selection strategy, the filtering is performed first, then the selection strategy is applied to the results of the filtering. Manual selection overrides both filtering and the selection strategy.

The Filter set option menu provides three preset filters, Druglike, Leadlike, and Coarse, a customizable filter (Custom), and an option to bypass filtering (None). If you modify one of the preset filters, the menu item switches to Custom.

Filtering is not available unless you have combinatorial docking results.

To use a predefined filter:

1. Choose a filter from the Filter set option menu.

To use a custom filter:

1. Choose a filter set from the Filter set option menu.

Choose Custom for a blank filter form, or choose one of the predefined filter sets to alter.

2. Enter the number of conditions that must be met for the molecule to be discarded in the Discard structure if N or more conditions are met text box.
3. Select the properties that you want to include in the filter set.
4. Enter the values in the text box that define the property range.

9.4 Results of Selection and Filtering

The selection and filtering job is run when you click OK or Apply in the Filter and Select dialog box. After the job finishes, the results of the selection and filtering process are displayed in the table in the upper part of the panel. The results—the current filter and selection settings and the results from these settings—constitute a library selection. The library size is displayed next to the Create Library button. The first four columns of the table are identical to the corresponding columns of the Reagent retention table from the Dock Library step. The columns are described in [Table 9.2](#).

Table 9.2. Description of the results table.

Column	Description
Name	Attachment position label
Reagent File	Name of the reagent file (minus the extension).
Single	Number of reagents from this file that survived the single-position docking.
Comb	Number of reagents from this file that survived the combinatorial docking. This number includes all reagents that contributed to any structure that docked well.
Library	Number of reagents selected for the library.
Best	Best GlideScore of the molecules used to select reagents at this position for the library. If combinatorial results are used in the selection, the best GlideScore is the same for all reagents.
Worst	Worst GlideScore of the molecules used to select reagents at this position for the library.

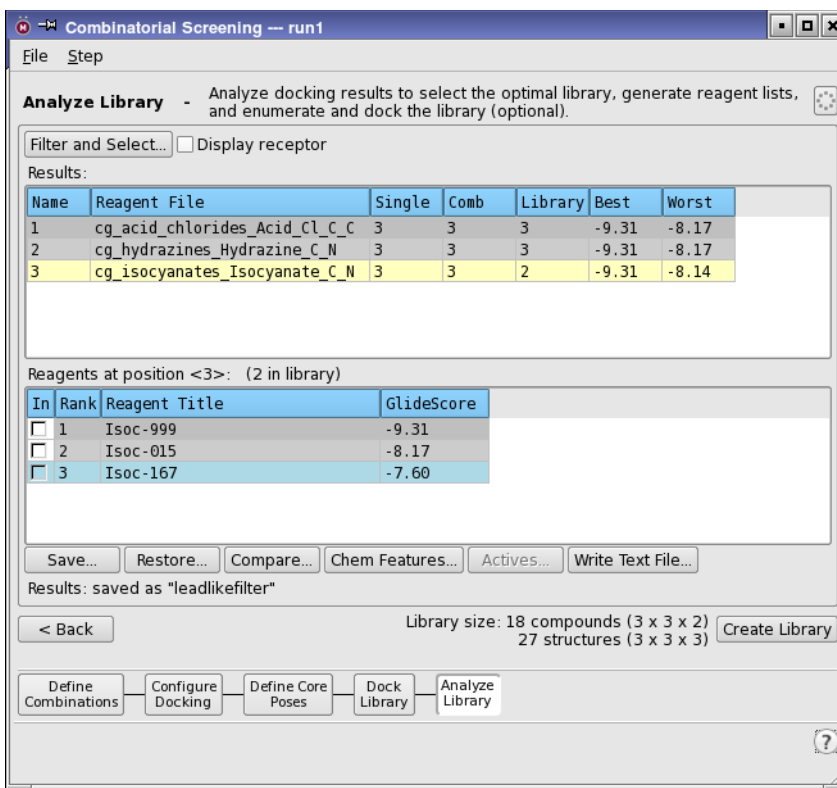


Figure 9.4. The Analyze Library step after selection.

When you select a row in the results table, the reagent list with results for that position is displayed in the Reagents at position *N* table. This table displays the top reagents for the position selected in the results table along with various properties (see Table 9.3). The list includes both the reagents that were chosen for the library at this position and a number of the best reagents that were not chosen. The reagents that were not chosen for the library are highlighted in blue.

In addition to examining the numerical results of the selection and filtering, you can view the best-scoring structure for each reagent in the Workspace. To do so, select the desired position in the results table, then click the In column of the Reagents at position *N* table for the reagent that you are interested in. The best-scoring singly substituted or fully substituted structure that contains the reagent is displayed. If you want to view the structure with the receptor, click Display Receptor. To remove the receptor, clear the Workspace.

If you decide that you want to include or exclude a reagent at a particular position, right-click in the row of the Reagents at position *N* table for the reagent and select Add to Include List or

Table 9.3. Description of the Reagents at position table.

Column	Description
In	Workspace inclusion status. You can use this column to display the best-scoring molecule for that reagent in the Workspace. However, manually included reagents that did not contribute to any docked molecule in the docking stage cannot be included in the Workspace.
Rank	Numerical rank of the reagent. The reagents are listed in order of rank; the rank is determined by the lowest GlideScore of any molecule that contains this reagent. A + sign indicates that the reagent was manually selected for inclusion in the focused library. Manually included reagents are listed at the top and are assigned the top ranks.
Reagent Title	Title of the reagent. This is the title chosen in the Reagent Preparation panel.
GlideScore	Best GlideScore for any compound in the library that used this reagent.
Boltz Pop	Boltzmann population for the reagent. Boltzmann population strategy only.
Delta G Rel	Effective relative GlideScore for this reagent treated as a free energy, derived from the Boltzmann populations. The best reagent has a Delta G Rel of zero. Boltzmann population strategy only.
Delta G Abs	Effective absolute GlideScore for this reagent treated as a free energy, derived from the Boltzmann populations. Boltzmann population strategy only.

Add to Exclude List from the shortcut menu that is displayed. Selecting a menu item opens the Filter and Select dialog box so that you can re-run the library selection job with the newly included or excluded reagent.

Running a new selection and filtering job overwrites the results of the current job (the current library selection). You can save the results in the run, by clicking **Save**. The **Save Results** dialog box is displayed, in which you can name the library selection. Once you have saved a library selection, you can restore it later as the current selection by clicking **Restore**, then selecting the library selection in the **Restore Results** dialog box and clicking **OK**. You can also delete a library selection in the **Restore Results** dialog box by selecting it and clicking **Delete**. The name of the current library selection is given immediately above the **Save** button—or if it is not saved, a message to indicate that has not been saved.

Saved library selections are stored within the run, and are not available outside the project. To save a library selection to an external file, click **Write Text File**, and navigate to the desired location in the file selector that is displayed. The contents of the file include the options used to run the library selection job, and the Reagents at position *N* table for each attachment position.

9.5 Comparing Results

In the process of selecting the optimal library, you may want to compare several library selections. To do so, these definitions must be saved in the run, using the Save button. To compare the saved library selections, click Compare, then in the Compare Results dialog box, select the library selections in the Select results to compare table. You can select and deselect items in this table with the usual click, shift-click and control-click actions. As you select the item, the Reagents at position *N* table for that library selection and the chosen position is displayed to the right of the list. To change the position for which results are displayed, choose the position from the Position menu. All the tables are updated with the results for the new position. You can compare up to six library selections.

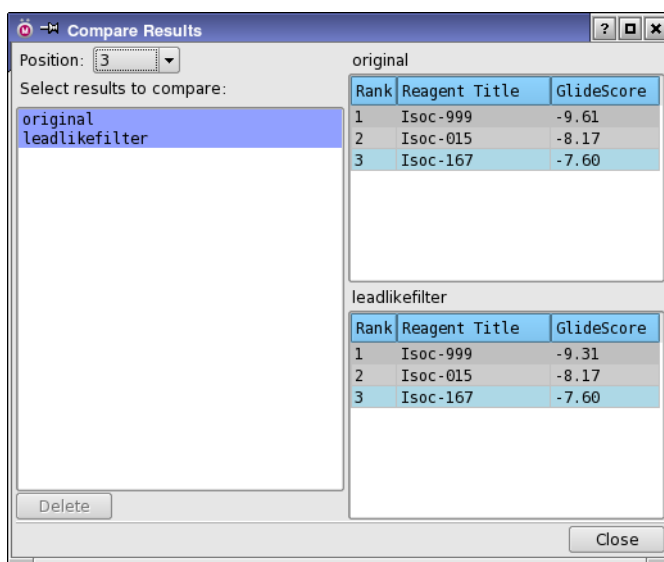


Figure 9.5. The Compare Results dialog box.

You can also delete library selections, by selecting them in the Select results to compare list and clicking Delete. The selections are deleted without further confirmation, so check your choices carefully.

9.6 Creating a Library

When you have decided on the library selection, you can proceed to create the library. The library is a Maestro file that contains structures generated from the reagent variants, subject to a GlideScore threshold. A reagent variant is discarded if the GlideScore of the best-scoring molecule to which it contributed is above the threshold. The threshold is the highest GlideScore of any molecule that contributed a reagent to the library. In addition, a plain text file is written that contains a summary of the library selection. This is the same file as is written with the Write Text File button. The library can be created with or without docking the results.

If the library is docked, the output Maestro file includes the receptor and the docked structures. The library can be docked with any of the Glide docking modes, but you must have a separate Glide license. If you choose Glide SP docking or Glide HTVS docking, the Post-docking section becomes available, and you can select options for post-docking minimization and applying strain correction terms. For more information on these options, see [Section 5.8.2](#) of the *Glide User Manual*.

If the library is enumerated only, you can choose to untangle and minimize the structures. When the structures are built, it is possible that the structures are “tangled”—for example, a chain might go through the center of a ring. In CombiGlide you can choose to untangle these structures, and then minimize them. Untangling and minimizing the structures takes longer, but the results are superior. Untangling is performed by default in the Dock Library step when the fully substituted structures are created.

To create the library:

1. Click Create Library.

The Create Library - Options dialog box opens.

2. Choose a mode for creating the library.

You can select Enumerate only, do not dock, or Enumerate and dock with.

3. Select options for the library creation mode.

4. Click OK.

A Start dialog box opens.

5. Select job options, and click Start.

Note: You cannot create a library for a command-line run that was imported with File → Import.

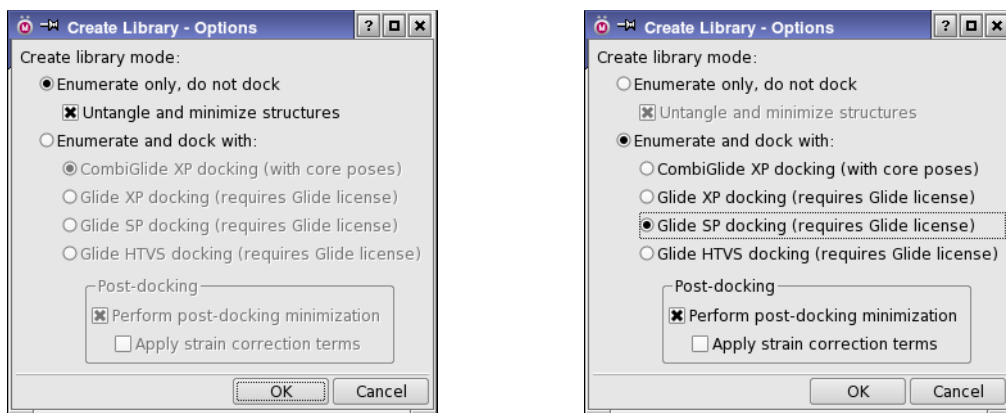


Figure 9.6. The Create Library - Options dialog box showing options for the two modes.

Combinatorial Screening: Analyzing the Library

When you have selected a small library of compounds that is likely to bind well, you may want more information about the library than is given by the GlideScore or the filtering. For example, the full library might include known actives, and you want to know how well the focused library recovers these actives. Or you might want to know what kinds of chemical features at each attachment position are found in the compounds that bind well. CombiGlide provides several tools for analyzing the results of a combinatorial docking run to obtain insight into what contributes to binding.

10.1 Chemical Feature Analysis

When selecting a small, high-ranking combinatorial library, you may be more interested in trends than in specific reagents. For example, you might like to know if hydrogen-bond donors are more likely than average to occur at a particular attachment point. These questions are addressed by the chemical feature analysis capabilities of CombiGlide.

Chemical feature analysis is done by classifying functional groups in side chains into chemical feature categories. The classification of the functional groups is performed automatically in the reagent preparation step, and the information is added to the reagent files. This classification is used later when a selection is performed using the Filter and Select dialog box. As part of the selection process, the chemical feature analysis is performed. The library must be one that was generated from a full combinatorial docking job: single-position results cannot be used.

When you have run a selection job, you can click **Chem Features** in the Analyze Library step to open the Chemical Features panel, which presents the results in graphical form. The panel and its display modes are described in subsections below.

If you want to apply chemical feature analysis to a CombiGlide 1.0 run, you can perform the classification on the reagent files with the command-line tool `cg_add_chem_features`, which is described in [Appendix B](#). Once you have converted the reagent files, you can go to the Analyze Library step to run the chemical feature analysis, without having to repeat the entire docking run. When you rerun the selection process from the Filter and Select dialog box, chemical feature results are added to the output.

The results are more likely to be meaningful when a small library is created in the Filter and Select run—say, a 1000-compound library, if the virtual space has size 1M. This is because the strongest binders, which will predominate in a small library, can be expected to most strongly

discriminate in favor of specific chemical features. On the other hand, a library that is too small will not provide a good statistical sampling of good binders.

10.1.1 Features and Types

Chemical features can be classified into one of the following categories:

- A—hydrogen bond acceptor
- D—hydrogen bond donor
- H—hydrophobe
- N—negatively charged center
- P—positively charged center
- R—aromatic

These categories reflect standard pharmacophore types, as used by Phase, for example. However, some of the chemical feature definitions used by CombiGlide differ from standard pharmacophore definitions.

- Classification as N and P is only done when the structure is actually in the ionic form. This is because CombiGlide docks the charged and uncharged structures separately.
- A side chain (the part of a reagent that remains after removing the reactive functional group that is discarded when it is bound to the core) may bear multiple features.
- A side chain is considered to be hydrophobic only if it is hydrophobic overall, not if it bears only a few connected nonpolar carbons. This definition is designed to weed out small hydrophobes in otherwise hydrophilic molecules. A structure is considered hydrophobic if any of the following conditions is met:
 - a. It has only hydrophobic features.
 - b. It has hydrophobic features containing five or more atoms. This includes most aromatic hydrophobes.
 - c. It has a QikProp-calculated logPo/w greater than some cutoff value (2.5 by default) and also contains a hydrophobe, however small. This condition prevents polar molecules containing a few hydrophobic carbon atoms from being considered hydrophobic. The QikProp calculation is performed by replacing the reactive functional group in the side chain with a methyl group.
- A single group in a side chain can belong to multiple feature categories. For example, carboxylates are acceptors (A) as well as negative centers (N), and a hydrophobic aromatic group is both hydrophobic (H) and aromatic (A).

We define a *type* as the list of features that a side chain has. For example, a side chain that has A, D and H features (one or more of each) will have type ADH. Because every side chain has

exactly one type, the sum of the fractions of side chains bearing each type is unity. This permits several analyses that are not available for features.

Feature and type analysis are both based on comparison between frequencies in the reagent collection used to define the full virtual library at a given position on the core (the initial collection) and the frequencies that appear at the same position in the selected library (the final collection).

The broader the distribution of types in the full library, the more discrimination it is theoretically possible to observe in the final collection. For example, if the side chains in the initial collection contain only hydrophobes, the selected library can contain only hydrophobes, and no enrichment would be visible. Thus, chemical feature analysis is not useful if the initial reagent collection is highly focused around chemical features and types known to be active at the site in question.

10.1.2 Displaying Feature and Type Analyses

Three analyses are presented in the Chemical Features panel, which can be selected with the options at the top of the panel:

- Side Chain Features
- Overall Type Selectivity
- Type Preferences

These analyses are described in the following subsections.

You can save the figures that are displayed in the panel in PostScript or GIF format, by making the relevant choice from the File menu. Each file contains only one figure—the one that is displayed—so you must display each figure in turn and save it to get a complete set of figures for all positions.

You can also export the data on which the figures are based either as a CSV file or a plain text file, again by making the relevant choice from the File menu.

10.1.2.1 Side Chain Features

The initial and final feature frequencies are displayed side by side. Error bars show the standard error, based on counting statistics. When the two frequencies differ by more than one standard error, the difference is significant. The results shown below indicate strong enrichment of acceptors, donors, and positive features for the position shown.

From this chart you can also obtain information on elimination of features, or on requirements for features. If the fractional appearance for the selected library is zero, the feature has been eliminated; if it is one, the feature is required.




Figure 10.1. Chemical Features panel, showing Side Chain Features chart.

To change the position, click Next Position or Previous Position. These buttons cycle through the attachment positions and display the chart for that position.

10.1.2.2 Type Selectivity

Entropy-based statistics can be used for fractional compositions that sum to unity, such as for the type data. We use a measure based on the Kullback-Liebler divergence to compute the selectivity of each position on the core. If a position has a selectivity of N , it is similar to saying that only one out of N types is selected at that position. This value would in fact apply if the types were uniformly distributed in both the initial and the final distribution, but the initial side-chain collection contained N times as many types as the final one.

In the following example, we can see that position 3 of this system has a selectivity of 7, which is much greater than the selectivity observed at the other positions, which is close to 1. Position 3 is the position for which the feature data were shown earlier. It is the high degree of discrimination in types that leads to the high selectivity.




Figure 10.2. Overall type selectivity chart.

Selectivity is always relative to the diversity of the reagent set used. If the input reagent set is already narrowly focused on the reagent types that bind, the observed overall selectivity will be low, even if the position is highly selective relative to all possible chemical types.

For positions with low selectivity, the type selectivity details can help determine whether the low selectivity is due to low diversity in the reagent set or to greater discrimination at other positions. A simple measure is obtained from the number of types that appear in a reagent list. More types means more diversity. The Type Selectivity Details chart, which you can display by clicking Show Type Details, shows the number of types in the input reagent collections in pale blue and the number of types in the selected library in bright blue. An example is given in [Figure 10.3](#).

In the example used for the figures, it can be seen that by this measure, the input reagent collections used at positions 1 and 2 are about as diverse than those used at position 3. Thus, using this measure of diversity, it could be concluded that the greater overall type selectivity exhibited by position 3 is intrinsic, rather than due to the use of a more diverse reagent collection at that position.

A better measure of diversity takes into account the non-uniform distribution of types at a position. For example, a two-type collection is more diverse if the types occur with equal frequency than if one type occurs 99% of the time and the other occurs 1% of the time. The latter situation is little different from having only a single type present. Such a measure that is based on the statistical entropy is also included in the type selectivity details. The entropy is converted into an effective number of types. These effective numbers are displayed in light green (for the input set) and dark green (for the selected library).




Figure 10.3. Type selectivity details chart.

For the example above, the effective number of types does not change much between the positions, so again it can be concluded that the high overall type selectivity exhibited at Position 3 is intrinsic, rather than artifactual.

10.1.2.3 Type Preferences

This chart visualizes how the distribution of chemical types changes between the initial and final collections. The upper bar in the display represents the initial side-chain collection and the lower bar represents the final side-chain collection. The width is normalized to 100% in both cases. Each vertical segment represents a chemical type, and each type has horizontal bars colored in for each feature that is part of the type; the width of each type represents its fractional appearance. For example, in the chart shown, the type farthest to the left, which we designate ADHPR, contains acceptors, donors, hydrophobes, positive centers and aromatics, but no negative feature. This type comprises only a few percent of the initial side-chain collection but about 50% of the side chains in the selected library. The third type, which may be designated AHR, is also strongly enriched in the selected library.

You can change the sort order by selecting one of the three options in the lower part of the panel: Alpha Sort sorts the types in alphabetical order of the type label; Input Reagents Fraction Sort and Selected Library Fraction Sort sort in order of fractional appearance in the respective collection.




Figure 10.4. Type preferences plot.

10.1.3 Using Chemical Feature Analysis Results

The results of a chemical feature analysis can be used in library refinement. Based on results with a diverse side-chain collection, a second-generation library can be designed that is enriched in the types and features shown to dominate at the various positions. Positions that exhibit low selectivity are less sensitive to the types and features that appear there, and a possible strategy is thus to use larger reagent collections (in proportion to selectivity values) at the more selective positions.

Two warnings, one pertaining to focused collections and one pertaining to diverse collections, must accompany this view:

- If the initial side-chain collection is already highly focused, it will not be possible to observe strong enrichment or high selectivity.
- If a diverse side-chain collection is too small, it might not be possible to find enough good binders at all positions to obtain meaningful results, due to inadequate sampling of chemical space. This will also lead to low enrichment and selectivity.

10.2 Active Compound Analysis

If you are interested in how well the focused library picks up known actives, you can use the active compound analysis capabilities to generate enrichment data and an enrichment plot. The analysis is done when you run a selection job from the Filter and Select dialog box.

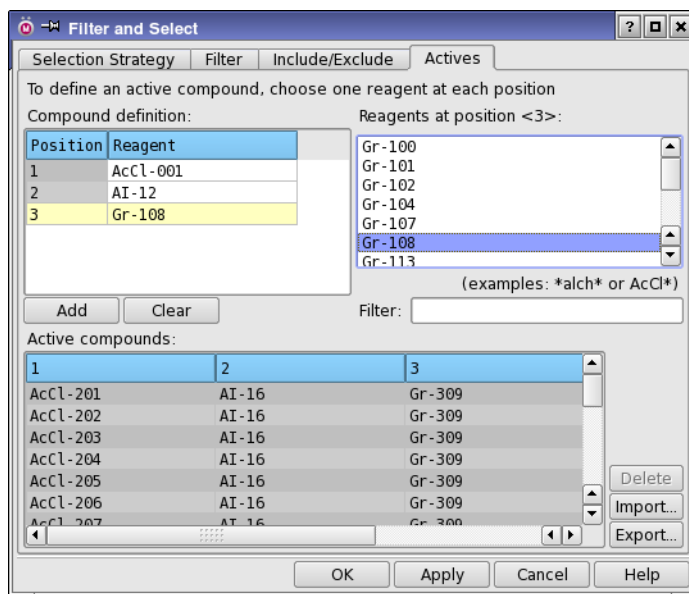


Figure 10.5. The Active tab of the Filter and Select dialog box.

The actives are specified in terms of a selection of the reagents at each position—they must, naturally, be members of the fully enumerated library in order to calculate an enrichment. The selection is made in the Active tab of the Filter and Select dialog box, as follows:

1. Choose a position in the Compound definition table.
2. Choose a reagent from the Reagents at position *N* list.

You can filter the reagents by entering text in the Filter text box, with * to represent an arbitrary string and ? to represent an arbitrary single character. You can also specify alternatives by listing them in braces { }, separated by commas. For example, 3-{F,Cl,Br}* would match any reagent that began with 3-F, 3-Cl, or 3-Br. All reagents whose titles contain the specified string are listed in the Reagents at position *N* list when you press ENTER.

You can also edit the Reagent column in the Compound definition table directly, but the name you enter must be one on the list. Entering text sets the filter in the Filter text box.

If you want to clear the Reagent column at the selected position, click Clear.

3. Repeat [Step 1](#) and [Step 2](#) for the remaining positions.




Figure 10.6. Actives dialog box.

4. Click Add.

A row is added to the Active compounds table. There is one column in this table for each position, which lists the reagent used at this position to define the active compound.

When you have finished defining the actives, you can run the library selection. Following the run, you can view information on the recovery of actives by clicking the Actives button in the Analyze Library step. This button opens the Actives dialog box, which shows the results in tabular form. The columns of the table are described in Table 10.1. From this dialog box, clicking Plot Enrichment Factors displays a standard enrichment plot, in the Enrichment Factors panel. The panel presents of a plot of the fraction of actives found against the fraction of the database screened, along with a plot of enrichment factors. You can plot these quantities against the library size instead by selecting Show Library Size.

This plot only extends to the range of the library selected. To see a full enrichment curve, set Max library size in the Filter and Select dialog box to a size exceeding that of the entire virtual library, then display the enrichment plot based on these results. This plot will not extend to the full virtual space, but it will cover the range accessible from the reagents used in the combinatorial docking (allnode) run.

An example of an enrichment plot is shown in Figure 10.7.

Table 10.1. Description of columns in the Actives found in selected library table.

Column	Description
Index	Index number of the active compound in the list of actives.
Compound	Identifier of the active compound, given in terms of the reagents used.
Level Found	Level in the selection procedure where this active compound was included in the library.
Lib Size Found	Size of the selected library when this active compound was found.
Lib Fraction	Fraction of the virtual library that was screened when this active compound was found. Equal to Lib Size Found divided by the total virtual library size (All reagents value in the Size of virtual libraries text area).
Fraction Active	Fraction of actives identified (including this compound) when this compound was identified in the selection procedure.
Enrichment	Enrichment factor, equal to the ratio of the fraction of actives identified (Fraction Active) to the library size ratio (Lib Fraction).

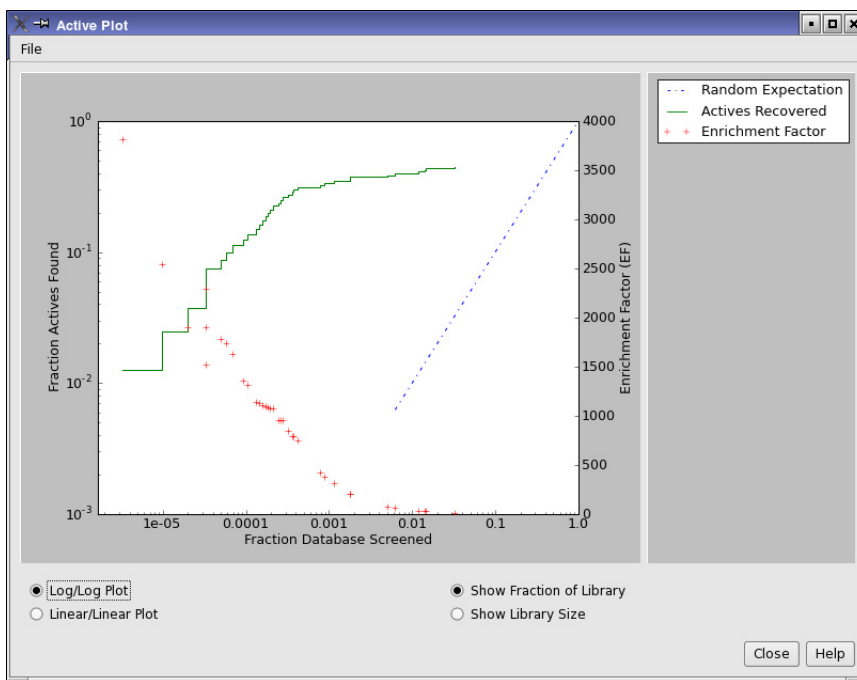


Figure 10.7. Enrichment factor plot.

Running Tasks from the Command Line

This chapter describes how to run various CombiGlide tasks from the command line. For information on command options, run the command with the `-h` option.

11.1 Reagent Preparation

Reagent preparation is implemented as a set of additions to LigPrep, which includes two scripts, `vcsprep` and `reagentprep`, and a command-line option to the `ligprep` command. The script `vcsprep` is run at the beginning of the LigPrep workflow, to ensure that the titles of the compounds are unique, and again at the end, to generate the statistics file and the SQLite database file that are used by CombiGlide. The script `reagentprep` is run on the structures produced by LigPrep, to add the labeling of the attachment bond. The concept of reagent preparation is described in detail in [Section 4.2 on page 25](#). You need not run either of these scripts.

To prepare a single reagent file for use in the combinatorial screening workflow or in combinatorial library enumeration, you can use the `ligprep` command as follows:

```
$SCHRODINGER/ligprep ligprep-options -vcs functional-group -omae jobname.bld
```

The *ligprep-options* are the usual LigPrep structure preparation options. The `-vcs` option specifies the name of the functional group to use. The names of the standard functional groups supplied with CombiGlide are the long names given in [Table 4.1 on page 29](#). You can also obtain a list of names with the following command:

```
$SCHRODINGER/utilities/reagentprep -list
```

Information on providing custom functional groups is given in [Appendix A](#). The output file must have a `.bld` extension, which indicates among other things that the file contains the information on the attachment bond (“grow bond”). The input file can be any file accepted by LigPrep.

You can also distribute the reagent preparation job. However, because there is no checking for duplicates between subjobs, the procedure used to ensure that titles are unique is to append the subjob number to each title, in addition to adding a suffix to duplicate titles. If this behavior is not acceptable, you should run the job serially.

In addition to the main output file, three structure files are produced by the `reagentprep` stage that contain structures that were rejected:

- *jobname_rp-mult.mae* contains structures that have multiple instances of the functional group specified that are not equivalent. Structures that have equivalent instances, such as the carboxyl groups in adipic acid for a carboxylic acid functional group, are accepted by default and one instance is labeled. Inequivalent instances, such as the carboxyl groups in glutamic acid, are rejected. If you want to discard such structures you can use the `-mult d` option to `reagentprep`, which you can pass through the `ligprep` command with the option `-W "rc,-mult,d"`. You can also include all instances of structures that have inequivalent instances with the `-mult a` option.
- *jobname_rp-nomatch.mae* contains structures that did not match the functional-group specification. If the input file should contain structures that all match, this file should be checked to determine why the structures did not match. If the input file is a master file of reagents for all functional groups, and reagent preparation is used to filter the structures, then this file merely contains the structures that did not match.
- *jobname_rp-ring.mae* contains structures for which the bond to be broken occurs in a ring. The Schrodinger build system cannot at this point handle such structures.

You should not use the output from a reagent preparation job as input to another job. If you do, the job will fail. If you want to label the reagents differently, you should start with the input file from the original job.

11.2 Combinatorial Library Enumeration

Enumeration of a combinatorial library can be done with the program `combgen`. The library can be constructed from a prepared core-containing structure, in which substitutions are made at the defined attachment points, or it can be constructed by linking several fragments in a linear sequence. The substitutions on a core-containing structure can also consist of several fragments.

The fragment libraries used to generate the combinatorial library can be taken from the standard Maestro builder (the default), or from reagent libraries prepared with `reagentprep`. If you want to use both, you must copy the libraries to a common location. Both the `.bld` and the `.frg` files must be copied, and the `.ini` file must be modified to include all desired libraries.

11.2.1 The `combgen` Command

The syntax of the `combgen` command is as follows:

```
$SCHRODINGER/combgen jobname [options]
```

11.2.2 The combgen Input File

The specification of the combinatorial library is made in the input file, which consists of a set of command lines. Each valid command line consists of a command name, followed by a colon (:), followed by a list of parameters specific to the command. The command names are case-insensitive. Comments can be added to the file after a # character.

There are four commands: `Core`, `Chain`, `Frag`, and `Vlib`.

The first command must be a `Core` command, and there must be only one `Core` command in the input file. The `Core` command can be followed by one or more `Chain` commands. Each `Chain` command should be followed by the appropriate number of `Frag` commands. `Chain` and `Frag` commands can be interspersed as long as each `Frag` command is preceded by the `Chain` command it refers to.

`Vlib` commands are optional. They can be used to specify and report on virtual fragment libraries, which are defined as subsets of the actual, full fragment libraries. Any `Vlib` commands that initiate or alter virtual fragment library fragment subsets must precede the first `Frag` command in the input. These `Vlib` commands are typically placed at the top of the file. Any `Vlib` commands that merely report on fragment subsets can occur after `Frag` commands, and are typically placed at the end of the input file.

The structure of a typical input file can be symbolically represented as follows:

```
Core
Vlib
...
Vlib
Chain
Frag
...
Frag
Chain
Frag
...
Frag
...
Vlib
...
Vlib
```

Annotated examples of input files are provided in `$SCHRODINGER/combiglide-vversion/samples`.

The commands with their syntax and use are described in the following subsections.

11.2.2.1 The Core Command

The Core command specifies the core-containing molecule, if one is to be used. The syntax is as follows:

Core: *filename*

The file must be in Maestro format, and the first structure from this file is taken as the core-containing molecule. If no file name is given, a coreless chain is assumed.

The structure in the file must include the grow name labels required to specify attachment points. These names specify the “from” atom (the atom from which the fragment is grown), and the “to atom” (the atom towards which the fragment is grown, and which is replaced, along with any atoms attached to it, when the fragment is added). The convention to use for these atoms is *nf* for the “from” atom and *nt* for the “to” atom, where *n* is the index of the attachment point, starting from 1. This is the convention that is applied when you label a core-containing molecule in Maestro. If you want to add these names manually, you should add them to the `s_m_grow_name` field in the atom block of the Maestro file.

11.2.2.2 The Chain Command

The Chain command specifies the chain to be attached at a given attachment point on the core-containing molecule. The chain can consist of multiple fragments, joined in a linear sequence.

Chain: *chain-name min-frags max-frags from-name to-name*

The fourth and fifth parameters apply only when a core has been specified. The parameters are described below.

<i>chain-name</i>	Name of the chain. This name is used by Frag commands to refer to the chain. The name may not contain any spaces, even if the name is quoted.
<i>min-frags</i>	Minimum number of fragments to build in this chain in exhaustive combinations. The lowest allowed value is 0 when there is a core, or 1 for a coreless chain.
<i>max-frags</i>	Maximum number of fragments to build in this chain in exhaustive combinations. The lowest allowed value is <i>min-frags</i> (or 1 if <i>min-frags</i> is 0). If <i>min-frags</i> = <i>max-frags</i> , all combinations will have the same number of fragments substituted.
<i>from-name</i>	Grow name of the “from” atom of the core molecule bond to break when attaching the first fragment of the chain.
<i>to-name</i>	Grow name of the “to” atom of the core molecule bond to break when attaching the first fragment of the chain.

When there is no core, there must be exactly one Chain command. In this case, the first fragment functions as the core, with a single attachment point defined by the forward grow direction. Even though there is only one chain, it must be given a name, to which subsequent Frag commands can refer.

The following example specifies a chain named ChainA, with 1 to 2 fragments added, starting at the attachment point on the core defined by grow names 1f and 1t:

```
Chain: ChainA 1 2 1f 1t
```

The following coreless chain specification leads to a combinatorial library in which every member consists of 4 fragments bonded together:

```
Chain: dummy 4 4
```

11.2.2.3 The Frag Command

The Frag command specifies the fragment library that is to be used to attach fragments to the specified chain at the specified fragment number. There must be a Frag command for each fragment, from fragment 1 to *max-frags*, of each specified chain. The syntax of the Frag command is as follows:

```
Frag chain-name frag-number lib-name
```

The parameters of the Frag command are specified below:

<i>chain-name</i>	Name of the chain to which this fragment will be attached. The name must match one of the names assigned in an earlier Chain command; the matching is case-insensitive.
<i>frag-number</i>	Fragment number (position) in the chain. The number can range from 1 to the value of <i>max-frags</i> for the specified chain. Fragment numbers increase moving away from the core.
<i>lib-name</i>	Name of the fragment library. All of the fragments in the named library are used in generating combinations. You can specify either an actual fragment library or a virtual fragment library (defined by a prior Vlib command). The library names of actual fragment libraries are those listed in the file specified with the -m command option (which by default is the file \$SCHRODINGER/mmshare-vversion/data/res/mmfrag.ini).

Fragment libraries that are used in a non-terminal location in the chain must have two grow bonds defined, a “from” bond and a “to” bond. The builder fragment libraries supplied with Maestro (in \$SCHRODINGER/mmshare-vversion/data/res) have two grow bonds defined.

The reagent libraries prepared with `reagentprep` have only one grow bond defined, and must be used in a terminal location in the chain.

11.2.2.4 The Vlib Command

Vlib commands are used to specify virtual fragment libraries, which are defined to be subsets of normal, full fragment libraries. A virtual fragment library can be specified in a `Frag` command for the purpose of limiting fragment selections in combinatorial generations.

The syntax of the Vlib command is as follows:

Vlib: *subcommand parameters*

There are five subcommands, whose syntax and use is described below.

Vlib: new *vlib-name fraglib initialization*

The new subcommand creates a new virtual fragment library specification, and initializes the virtual library from an actual library. Subsequent `add` or `delete` subcommands define the subset of this library that is used in the virtual library.

<i>vlib-name</i>	The name of the virtual fragment library. This name must not match an actual fragment library or another virtual fragment library.
<i>fraglib</i>	The name of the actual fragment library of which this virtual fragment library is a subset.
<i>initialization</i>	Keyword that specifies how to initialize the virtual fragment library. Valid values are: all full—Initialize the virtual library with all members of the full library. none empty—The virtual library is initially empty.

Vlib: add *vlib-name frag-name-list*

The add subcommand adds a specified list of fragments to the given virtual fragment library. The list is a blank-separated list of fragment names. Names of fragments that have already been added are ignored. You can include as many add subcommands as you want. This subcommand is normally used for a virtual fragment library that is initially empty.

Vlib: remove *vlib-name frag-name-list*

The remove subcommand removes a specified list of fragments from the given virtual fragment library. The list is a blank-separated list of fragment names. Names of fragments that have already been removed are ignored. You can include as many remove

subcommands as you want. This subcommand is normally used for a virtual fragment library that is initially full.

Vlib: list *vlib-name*

The `list` subcommand displays the list of fragment names defined in the specified virtual fragment library. This subcommand must follow the `new` subcommand that defines the virtual fragment library. To generate a list of fragment names in a library, you can use this command immediately after a `Vlib: new` command with initialization to `all`.

Vlib: count *vlib-name*

The `count` subcommand displays the number of fragment names defined in the specified virtual fragment library. This subcommand must follow the `new` subcommand that defines the virtual fragment library.

All uses of the `Vlib` subcommands `new`, `add`, and `remove` must occur before the first `Frag` command. Uses of `list` or `count` can occur after the `Frag` commands.

When virtual fragment libraries are created, some temporary files are produced. These files are stored by default in `$HOME/.schrodinger/vlib_tmp`. You can specify the location for these files by setting the environment variable `SCHRODINGER_VLIB_TMPDIR`.

11.3 Library Selection

The selection of a full combinatorial library starting with a docked selection from a much larger virtual combinatorial space can be performed with the utility `libselector`. A full description of library selection is given in [Chapter 9](#).

The `libselector` utility produces a log file and also a file containing reagents comprising the selected library. `Libselector`'s input consists of $N+2$ SQLite database files, where N is the number of substitutable positions on the core in the library definition. For each position, there is an SQLite database that is produced by the reagent preparation procedure along with the fragment files *filename*.bld; these are named *filename*.sqlite. If the structural output of the combinatorial docking job is *jobname*.mae, then the `CombiGlide` procedure should also produce *jobname*.sqlite and *jobname*-comb.sqlite, both of which `libselector` requires.

11.3.1 The libselector Command

The syntax of the `libselector` command is as follows:

```
$SCHRODINGER/utilities/libselector -s strategy -S strategy_options  
-j jobname [options]
```

11.3.2 The QikProp Filter File

The custom QikProp filter file consists of a set of specification lines with one of the following forms

```
property-name min max
property-name op value
```

where *property-name* is the name of the QikProp property as it appears in a Maestro file, *min* and *max* are the minimum and maximum values for the property value, *op* is one of *gt* (greater than) or *lt* (less than), and *value* is a threshold. If the property value for a structure meets the specification, or the structure does not have the property, the structure passes the filter. The number of filter failures needed to reject a structure is given by a specification line in the file using the special property *i_qp_#rejections*. Comment lines may be included by beginning them with a # character.

A sample filter file follows:

```
i_qp_#rejections gt 1
r_qp_mol_MW      150 650
r_qp_QPlogPo/w   -3 6
r_qp_PSA         gt 174.9999
```

11.3.3 Include and Exclude Files

The include and exclude files have a common format. The file consists of one or more lines that define a position; each such line is followed by one or more lines defining a reagent. The position definition line has the following format:

```
position_name: name
```

where *name* is the name of the position. The reagent definition lines contain the name of the reagent. Comments can be added on lines beginning with a # character. Blank lines are ignored.

11.3.4 Actives File

Actives files are specified with either the *-a* option or the *-A* option. These files provide a list of specifications of which compounds in the library are known active.

For an active file specified by *-a filename*, the specification of an active consists of a list of comma-separated integers. Each integer is the index of the reagent at the given position (starting with 1), in the corresponding *.bld* files used to define the library. Since the *.bld* files could contain multiple modifications of any reagent, the integers are necessarily specifying specific modifications. However, all modifications of a reagent are considered equivalent

for the purpose of searching for actives. That is, the specification is understood to refer to the underlying reagent, not the specific modification.

For an active file specified by -A, the specification of an active consists of a list of reagent names, one for each position. Each reagent name is given in square brackets, so the following constitutes a valid example:

```
[3,4,5-trimethoxy-beta-phenylethyl amine] [acetic acid] [LMDR 426]
```

Within a pair of square brackets delimiting a reagent name, leading and trailing whitespace is ignored, but embedded whitespace is significant. Matching square brackets within a reagent name are accepted; unmatched embedded square brackets must be escaped with a backslash, thus: \[. In the two lines given below, the first specifies reagent name g whereas the second specifies reagent name [g], where the white space is significant:

```
[ g ]
[ [ g ] ]
```

11.3.5 The Reagent File

The main purpose of this file is to provide lists of reagents that the program selected for each combinatorial position. In addition, it provides information about the overall run and the reagents, and also gives a list of “runners up” at each position: reagents that might be selected if the maximum library size were increased.

The file has a header section, which contains information about the run as a whole, followed by sections containing information about each grow position. Comments are included on lines that begin with a # character. Many lines are in the form of name-value pairs, where a name is a string followed by a colon. The name is followed by optional whitespace then by the value.

In the following description, the exact names used in the file are given only when they are not sufficiently descriptive to be obvious when seen.

11.3.5.1 Header Section

This section always includes invocation: and args: lines, giving the program name and arguments as it was executed. It also includes values for various library sizes:

- lib_size: size of the library selected by libselector
- ext_size: the maximum library size buildable from the output of CombiGlide docking
- bld_size: the size of the original virtual space.

These sizes are in terms of reagents, not reagent modifications; for example, acetic acid counts only once, even if it occurs at some position in the selected combinatorial library in both ionized forms.

If QikProp filters were in use, the filter values are given on `qp_filter:` lines.

If `-a` or `-A` was used to specify a list of known actives, statistics for the known actives are given:

- `num_actives_unique`: number of unique active combinations given in the input
- `num_actives_found`: specifies the number found in the selected library

For each active compound, its position in the input list (starting with 0) is given, followed by its composition, followed by a vector of values pertaining to it and when it was found when building up the library. Values include the level found, the size of the library when it was found, what fraction of the original virtual library this constitutes, what fraction of the specified actives was found when this one was found, and the enrichment factor (the ratio of the fraction of actives found to fraction of the original virtual library comprising the libselector selection when the active was found).

The “level found” corresponds to the step the active was first found at as the library was being built. For the `best` strategy, this is the index of the GlideScore-sorted CombiGlide docked structure where the active was first found, omitting structures skipped for reasons of QikProp filtering, user-specified exclusion, or preemptive strategy. For the `boltz` strategy, it is the level at which the reagent was found when traversing the list of reagents in order of decreasing Boltzmann fraction.

11.3.5.2 Position Sections

Each a position section starts with a `position_name:` line, then a `reagent_filename:` line, which specifies the name of the `.bld` file used to supply reagents at this position. A list of user-excluded reagent names at this position, if any, comes next.

Following this, the file gives the best and worst GlideScores of structures contributing to the reagent list that follows, and the `number_selected:`, which is the number of modifications of reagents selected at this position (where a modification could be an ionization state, a tautomerization state, or a stereoisomeric state).

Then come a set of lines, one for each selected reagent modification. From the point of view of synthesis, only the reagents matter, but from the point of view of later enumerating the selected virtual library, it is important to list all high-scoring modifications.

For each reagent selected, a vector of values is given that depends, to some extent, on the strategy used. The level at which each reagent was found is given (see discussion above concerning the known actives in the header section); this is 0 for a user-included reagent that did not appear in the CombiGlide docking output. The structure in which each reagent first appeared in the docking output is given, as well as its index in the `.bld` file. These values pertain to this modification of the reagent.

For the `boltz` strategy, the Boltzmann population and the computed `Delta_G` of each reagent is given, relative to the most frequent reagent (in the Boltzmann sense) at its position.

For each reagent there is a 1 in the `b_first` column if it is the first modification of that reagent to appear in this position of the output, and a 0 otherwise. Thus, by counting the 1's, you can count the reagents, rather than the modifications.

A 1 is placed in the `b_include` column to indicate that the reagent was specified on the user's list of reagents always to include at that position. Finally, the reagent name is reported last on each line.

Following the `number_selected`: lines of selected reagents, a few runners-up are listed, as described earlier.

Creating Custom Functional Groups and Custom Minimal Capping Groups

This appendix describes the procedures for adding custom functional groups or modifying existing functional groups.

The definitions of the functional groups are stored in two files in the CombiGlide distribution, `reagentprep.ini`, which is a plain text file, and `reagentprep.mae`, which is a structure file in Maestro format. These files may be found in the directory `$SCHRODINGER/combiglide-vversion/data`. The functional group definitions include the following:

- **A long name.** This name is used to identify the functional group, and is used for its title. The long name is displayed in a tooltip when you pause the pointer over the reagent button in the Reagent Preparation panel. The name can be chosen any way you wish, but must not contain blanks. The long names for the supplied functional groups are constructed by adding the atom that is kept and the atom that is lost to the functional group name, separated by underscores, for example `Acid_Chloride_C_C`.
- **A short name.** This is the name that is displayed on the reagent button in the Reagent Preparation panel. The short names for the supplied functional groups are constructed by adding the atom that is kept and the atom that is lost to the shortened functional group name, separated by underscores, for example `Acid_Cl_C_C`. Because the reagent button displays the short name for custom functional groups, it is advisable to ensure that you can identify both the functional group and the bond that is broken in the name. Adopting the convention used for the supplied functional groups is one way to do this.
- **One or more SMARTS patterns that define the group.** The SMARTS patterns must be contiguous, without any spaces. At the top of `reagentprep.ini` is a description of the syntax with which you can use multiple SMARTS patterns to define the functional group.
- **A grow bond definition.** The grow bond defines the bond to be broken when the side chain from the reagent is added to the core. The bond is defined in terms of two atoms, the atom that is kept, and the atom that is deleted. This bond must be a single bond and must not be internal to a ring.
- **A structure that defines the minimal capping group.** This is the group that is added at the attachment position to define the minimally capped core, and in the construction of the singly substituted core molecules. The structure should consist of the desired functional group with the minimal capping group attached. For example, if you wanted to define an acid chloride with a methyl group as a minimal capping group, you would choose acetyl chloride as the structure. The structure should be 3D and minimized.

To add a new functional group, you must edit the file `reagentprep.ini` and add one or more lines that include the long name, the SMARTS pattern, and the grow bond definition, and you must add a structure to the file `reagentprep.mae`, with the long name, the short name, and the grow bond defined.

To change the capping group for an existing functional group, you must modify the structure in the file `reagentprep.mae`.

You can store the modified files in the directory `$HOME/.schrodinger/combiglide` or the directory from which you submit your reagent preparation job. When a reagent preparation job is run, CombiGlide looks first in the job submission directory for these files, then in `$HOME/.schrodinger/combiglide`, then in `$SCHRODINGER/combiglide-vversion/data`. If you want to customize these files just for a particular job, you should store them in the job submission directory. If you want to use the custom functional groups for all your jobs, you should store the files in `$HOME/.schrodinger/combiglide`.

To add a functional group definition to `reagentprep.ini`:

1. Copy the file `reagentprep.ini` to the desired location.
2. Open `reagentprep.ini` in a text editor.
3. Insert a line into the file with the following format, followed by a blank line:

```
AddGN longname SMARTSpattern kept-atom:rpc1 lost-atom:rpc2
```

An example for the acid chloride functional group is:

```
AddGN Acid_Chloride_C_Cl [CX3,c] (=O) (Cl) [#6] 1:rpc1 3:rpc2
```

The variables that you need to replace have the following meanings:

<i>longname</i>	The long name of the functional group (see above).
<i>SMARTSpattern</i>	SMARTS pattern that defines the functional group.
<i>kept-atom</i>	The index of the atom in the grow bond that is kept when the bond is broken and the side chain is added to the core. The index is the index in the SMARTS string.
<i>lost-atom</i>	The index of the atom in the grow bond that is lost when the bond is broken and the side chain is added to the core. The index is the index in the SMARTS string.

`rpc1` and `rpc2` are the *grow names* for the specified atoms, which are used to identify the kept atom and the lost atom when the side chain is “grown” onto the core.

4. Save and close the file.

To add a structure for the functional group to reagentprep.mae:

1. Copy the file `reagentprep.mae` to the desired location.
2. Start Maestro and open the Project Table panel (click the toolbar button shown below).



3. Add the desired structure to the Project Table.

You can import the structure using the Import panel, or you can build it, or duplicate one of the structures in the Project Table and modify it. It might be useful to import the structures from `reagentprep.mae` to duplicate and modify.

The structure must consist of the functional group and the minimal capping group only. The structure must be a 3D minimized structure. If it is not, you can convert and minimize it with LigPrep, or choose Clean Up Geometry from the Edit menu in the main window to minimize it with the UFF minimizer.

If you build the structure from fragments, a project entry is created automatically. If you use other means, you may have a scratch entry, and you must add it to the Project Table, by clicking the Create Entry button on the Workspace toolbar.



The title must be the long name of the functional group. You can also add the structure to the entry group that was imported, but it is not necessary to do so.

4. Enter the long name and the short name in the appropriate columns of the Project Table for the new structure.
5. Enter the long name as the title of the new structure, if you haven't done so.

For both of these steps, you should ensure that the long name is the same as in the file `reagentprep.ini`.

Now you define the grow bond.

6. Ensure that the structure is displayed in the Workspace.
7. In the main window, choose **Workspace > Atom Labels**.

The Atom Labels panel opens.

8. In the Composition folder, select Grow Name and deselect all other labels.

9. Click All in the Label Atoms section.

The atoms are now labeled with their grow names, if they have one.

10. Choose Edit > Build > Atom Properties.

The Build panel opens at the Atom Properties tab.

11. Choose Grow Name from the Property option menu.

12. Ensure that Atoms is chosen from the Pick option menu in the Apply grow name section.

13. Type `rpc1` into the Grow Name text box, and click on the atom in the Workspace that is to be kept when the bond is broken.

14. Type `rpc2` into the Grow Name text box, and click on the atom in the Workspace that is to be lost when the bond is broken.

If you make a mistake with any of these assignments, just type the correct name in and pick the atom again. These two atoms must correspond to the atoms in the SMARTS pattern whose indices you entered in `reagentprep.ini`.

To delete a grow name, enter two successive quotation marks "" (a null string) into the Grow Name text box.

Finally, you add the new structure to the file.

15. Choose Export Structures from the Project menu in the main window.

The Export dialog box opens.

16. Navigate to your copy of `reagentprep.mae` and select it.

17. Under the File text box, select Append.

18. Under Structure sources to be exported, select Workspace.

19. Ensure that Export all entries to a single file is selected from the Files option menu.

20. Click Export.

The entry displayed in the Workspace, which is your new structure, is appended to `reagentprep.mae`

If you want to add multiple functional groups, you can repeat the relevant steps above for other structures. When you come to adding the new structures, select the entries in the Project Table, and select Selected entries instead of Workspace under Structure sources to be exported in the Export dialog box ([Step 18](#)).

To modify a minimal capping group:

1. Copy the file `reagentprep.mae` to the desired location.
2. Start Maestro and open the Project Table panel (click the toolbar button shown below).



If you start Maestro in the directory to which you copied `reagentprep.mae`, no navigation is required when you import or export the file.

3. Import `reagentprep.mae`.



4. Include the structure for the reagent you want to modify in the Workspace.
5. Modify the structure.

You can use the Build panel to modify or add to the structure.

6. Minimize the structure.

The preferred treatment is to use MacroModel with the OPLS_2005 force field. Minimization is important because the docking process does not change the bond lengths and bond angles in the single-position docking step.

7. Select all the structures that were imported.

This should be all the structures in the Project Table, if you started with a new Maestro session.

8. Choose Export Structures from the Project menu in the main window.

The Export dialog box opens.

9. Navigate to your copy of `reagentprep.mae` and select it.
10. Under the File text box, deselect Append.
11. Under Structure sources to be exported, select Selected entries.
12. Ensure that Export all entries to a single file is selected from the Files option menu.
13. Click Export.

A dialog box is displayed to confirm the overwriting of `reagentprep.mae`

14. Click Yes.

The file `reagentprep.mae` is overwritten with the modified structure set.

To replace the existing structure for a capping group with another structure:

1. Follow the instructions for adding a structure above, from [Step 1](#) through [Step 14](#).

When you enter the title, long name, and short name, use the values for the entry containing the structure you want to replace.

2. Delete the entry that you want to replace.

The new entry has all the required information in it, and the order of the entries does not matter. If you want to move the new entry to the position that the deleted entry was in, select the entry and drag it to the desired location, or choose **Move to Row** from the **Entry** menu.

3. Select all the structures that were imported and the new structure.

If you started with a new Maestro session, this should be all structures in the **Project Table**.

4. Follow the instructions for modifying a minimal capping group, from [Step 8](#) through [Step 14](#).

Command-Line Tools

The CombiGlide distribution includes a number of utilities, which are installed into `$SCHRODINGER/utilities`, and are described below. For information on command options, run the command with the `-h` option.

B.1 `cg_add_chem_features`

This utility adds chemical features to CombiGlide reagent files that were originally prepared without them, so that chemical feature analysis can be performed on an already-completed CombiGlide run. Starting with Schrodinger Suite 2007, chemical features are added to all newly prepared reagent files automatically. Reagent files are identified by a `.bld` suffix. Along with this file there are also an `.sqlite` and a `_vpost-stats.txt` file, with the same base name. These three files are required by CombiGlide when adding a reagent.

The syntax is as follows:

```
cg_add_chem_features [-h] reagent-file [reagent-file2 ... ]
```

The reagent file name can be given with or without the `.bld` suffix. This file must include a block at the top containing an `s_rp_functionalgroup` property, which will always be the case if the file was created by means of the CombiGlide reagent preparation procedure.

The utility writes a new `.bld`, `.sqlite` and `_vpost-stats.txt` file for each input file, with the same names as the input files. The new files are always created in the current directory. If you want to avoid overwriting the old files, change to the directory in which you want to write the new files, and specify the input files using relative or absolute paths.

If the utility is run from a directory containing the old files, an attempt will first be made to back up the old files by appending `-bak` to their names, so `myfile.bld` is copied to `myfile.bld-bak`, and likewise for the `.sqlite` and `_vpost-stats.txt` files. If both the `myfile.bld` and the `myfile.bld-bak` files exist, the utility assumes that a previous run failed, and uses the `-bak` files as input and overwrites the original files.

To add chemical features to reagent files used in earlier runs of CombiGlide, you need to run the program on the copies of the `.bld` files that were stored when the run was done. If the run was done from a Maestro project `project.prj`, the files are in directories named `project.prj/.mmproj-admin/vcs.N` and `project.prj/.mmproj-admin/vcs.N/combiglide`, where `N` is an integer. You should process both sets of `.bld` files.

Finally, some Schrödinger programs require `.mae` suffixes. To accommodate these programs, a `myfile.mae` symbolic link is temporarily made to the old reagent file `myfile.bld`, and the link is removed when it is no longer needed. If `myfile.mae` already exists when the utility is executed, the utility fails.

B.2 chem_features

This utility lists the chemical feature types present in the input structures. Several definitions of feature types can be used: the CombiGlide definition, the Phase definition, or a user-specified definition. See [Section 10.1.1 on page 74](#) for a description of the standard feature definitions used by CombiGlide.

The command syntax is as follows:

```
chem_features [-h] [-v level] [-r] [-R recapped-file] [-p]
               [-f {P|C|filename}] [-q logP-cutoff|-Q] [-o output-file] input-file
```

In the absence of `-h` or `-v` options or errors, there is no text output.

Path lengths are computed as follows:

- For each atom in each pharmacophore, the minimum number of bonds between each atom in the pharmacophore and the capping atom is computed.
- The maximum of these values for the feature is the maximum path for that feature.
- The maximum of these values for all features of the same feature type (e.g., hydrophobes) is the value reported for each feature type.

If paths are computed then, for each feature type, properties `i_chemfeatures_type_atom` and `i_chemfeatures_type_dist` are computed. The `_dist` property contains the maximum path for that feature type and the `_atom` property contains an atom number exhibiting that distance. The possible values of `type` are hydrophobe, donor, acceptor, positive, negative, and aromatic.

If a feature of type `type` is not present, then the `_atom` and `_dist` properties for that type are given values of 0.

If there is no capping atom but there is an atom with growname `rpc2`, it is used to compute paths.

If the type is present but there is no capping atom, the `_atom` feature is given a value of -1 and the `_dist` feature is given a value of 0. Thus, if `-p` is specified, the `_atom` property for any feature type is zero if the feature is absent and non-zero if it is present.

Examples:

List default chemical features for the structures in `input.mae`:

```
$SCHRODINGER/utilities/chem_features -v 1 input.mae
```

List Phase pharmacophoric features for the structures in `input.mae`:

```
$SCHRODINGER/utilities/chem_features -v 1 -f P input.mae
```

Normal CombiGlide use, during reagent preparation, replacing the reactive side chain with $-\text{CH}_3$ and producing file `output.bld` containing the properties representing the properties in `input.bld`:

```
chem_features -r -p -q 2.5 -o output.bld input.bld
```

Verbose CombiGlide use, producing a file containing the capped structures as well as verbose output concerning operation of the program):

```
chem_features -v 2 -R capped.bld -p -q 2.5 -o output.bld input.bld
```

B.3 cg_combine_bld

This utility combines a set of reagent (`.bld`) files that were created by the Reagent Preparation procedure so that the combined file can be used as a reagent file at a single position in combinatorial enumeration or screening. Structures whose R-groups are identical to ones in previous structures processed are discarded.

The utility uses the functional group definition found in the first input file specified as the functional group definition in the output file. The functional group specification has the effect of defining the minimal capping group in the Combinatorial Screening workflow. However, since minimal capping groups are not used by default, the order of specification of input files does not matter unless minimal capping groups are explicitly used.

The utility can also convert a single fragment file in Maestro format as written by the R-Group Analysis panel into a reagent file. See *R-Group Analysis* for details on creating the fragment files.

The command syntax is as follows:

```
$SCHRODINGER/utilities/cg_combine_bld [options] input-list output
```

where *input-list* is a list of the reagent file stems (file names without the `.bld` extension) or a single Maestro file stem; *output* is the file stem of the output reagent file. The output consists of the `output.bld` reagent file, `output_vpost-stats.txt`, and `output.sqlite` files, and if

specified, *output.frg*. The last three files allow the output *.bld* file to be used as a reagent file in combinatorial screening.

A typical use for an R-group fragment file is:

```
$SCHRODINGER/utilities/cg_combine_bld -m -F -g Me_rga my_rga reagent1
```

which converts the Maestro file *my_rga.mae* into a reagent file *reagent1.bld*, and generates the *reagent1.frg*, *reagent1.sqlite*, and *reagent1_vpost-stats.txt* files. The functional group name is *Me_rga*.

Getting Help

Information about Schrödinger software is available in two main places:

- The `docs` folder (directory) of your software installation, which contains HTML and PDF documentation. Index pages are available in this folder.
- The Schrödinger web site, <http://www.schrodinger.com/>, In particular, you can use the Knowledge Base, <http://www.schrodinger.com/kb>, to find current information on a range of topics, and the Known Issues page, <http://www.schrodinger.com/knownissues>, to find information on software issues.

Finding Information in Maestro

Maestro provides access to nearly all the information available on Schrödinger software.

To get information:

- Pause the pointer over a GUI feature (button, menu item, menu, ...). In the main window, information is displayed in the Auto-Help text box, which is located at the foot of the main window, or in a tooltip. In other panels, information is displayed in a tooltip.

If the tooltip does not appear within a second, check that Show tooltips is selected under General → Appearance in the Preferences panel, which you can open with CTRL+, (⌘,). Not all features have tooltips.

- Click the Help button in the lower right corner of a panel or press F1, for information about a panel or the tab that is displayed in a panel. The help topic is displayed in the Help panel. The button may have text or an icon:



- Choose Help → Online Help or press CTRL+H (⌘H) to open the default help topic.
- When help is displayed in the Help panel, use the navigation links in the help topic or search the help.
- Choose Help → Documentation Index, to open a page that has links to all the documents. Click a link to open the document.

- Choose Help → Search Manuals to search the manuals. The search tab in Adobe Reader opens, and you can search across all the PDF documents. You must have Adobe Reader installed to use this feature.

For information on:

- Problems and solutions: choose Help → Knowledge Base or Help → Known Issues → *product*.
- New software features: choose Help → New Features.
- Python scripting: choose Help → Python Module Overview.
- Utility programs: choose Help → About Utilities.
- Keyboard shortcuts: choose Help → Keyboard Shortcuts.
- Installation and licensing: see the *Installation Guide*.
- Running and managing jobs: see the *Job Control Guide*.
- Using Maestro: see the *Maestro User Manual*.
- Maestro commands: see the *Maestro Command Reference Manual*.

Contacting Technical Support

If you have questions that are not answered from any of the above sources, contact Schrödinger using the information below.

Web: <http://www.schrodinger.com/supportcenter>
E-mail: help@schrodinger.com
Mail: Schrödinger, 101 SW Main Street, Suite 1300, Portland, OR 97204
Phone: +1 888 891-4701 (USA, 8am – 8pm Eastern Time)
+49 621 438-55173 (Europe, 9am – 5pm Central European Time)
Fax: +1 503 299-4532 (USA, Portland office)
FTP: <ftp://ftp.schrodinger.com>

Generally, using the web form is best because you can add machine output and upload files, if necessary. You will need to include the following information:

- All relevant user input and machine output
- CombiGlide purchaser (company, research institution, or individual)
- Primary CombiGlide user
- Installation, licensing, and machine information as described below.

Gathering Information for Technical Support

The instructions below describe how to gather the required machine, licensing, and installation information, and any other job-related or failure-related information, to send to technical support. Where the instructions depend on the profile used for Maestro, the profile is indicated.

For general enquiries or problems:

1. Open the Diagnostics panel.
 - **Maestro:** Help → Diagnostics
 - **Windows:** Start → All Programs → Schrodinger-2015-2 → Diagnostics
 - **Mac:** Applications → Schrodinger2015-2 → Diagnostics
 - **Command line:** \$SCHRODINGER/diagnostics

2. When the diagnostics have run, click Technical Support.

A dialog box opens, with instructions. You can highlight and copy the name of the file.

3. Upload the file specified in the dialog box to the support web form.

If you have already submitted a support request, use the upload link in the email response from Schrödinger to upload the file. If you need to submit a new request, you can upload the file when you fill in the form.

If your job failed:

1. Open the Monitor panel, using the instructions for your profile as given below:

- **Maestro/Jaguar/Elements:** Tasks → Monitor Jobs
- **BioLuminate/MaterialsScience:** Tasks → Job Monitor

2. Select the failed job in the table, and click Postmortem.

The Postmortem panel opens.

3. If your data is not sensitive and you can send it, select Include structures and deselect Automatically obfuscate path names.
4. Click Create.

An archive file is created, and an information dialog box with the name and location of the file opens. You can highlight and copy the name of the file.

5. Upload the file specified in the dialog box to the support web form.

If you have already submitted a support request, use the upload link in the email response from Schrödinger to upload the file. If you need to submit a new request, you can upload the file when you fill in the form.

6. Copy and paste any log messages from the window used to start the interface or the job into the web form (or an e-mail message), or attach them as a file.

- **Windows:** Right-click in the window and choose **Select All**, then press **ENTER** to copy the text.
- **Mac:** Start the **Console** application (**Applications** → **Utilities**), filter on the application that you used to start the job (**Maestro**, **BioLuminate**, **Elements**), copy the text.

If Maestro failed:

1. Open the **Diagnostics** panel.

- **Windows:** **Start** → **All Programs** → **Schrodinger-2015-2** → **Diagnostics**
- **Mac:** **Applications** → **SchrodingerSuite2015-2** → **Diagnostics**
- **Linux/command line:** `$SCHRODINGER/diagnostics`

2. When the diagnostics have run, click **Technical Support**.

A dialog box opens, with instructions. You can highlight and copy the name of the file.

3. Upload the file specified in the dialog box to the support web form.

If you have already submitted a support request, use the upload link in the email response from Schrödinger to upload the file. If you need to submit a new request, you can upload the file when you fill in the form.

4. Upload the error files to the support web form.

The files should be in the following location:

- **Windows:** `%LOCALAPPDATA%\Schrodinger\appcrash`
(Choose **Start** → **Run** and paste this location into the **Open** text box.)
Attach `maestro_error_pid.txt` and `maestro.exe_pid_timestamp.dmp`.
- **Mac:** `$HOME/Library/Logs/CrashReporter`
(Go → **Home** → **Library** → **Logs** → **CrashReporter**)
Attach `maestro_error_pid.txt` and `maestro_timestamp_machinename.crash`.
- **Linux:** `$HOME/.schrodinger/appcrash`
Attach `maestro_error_pid.txt` and `crash_report_timestamp_pid.txt`.

If a Maestro panel failed to open:

1. Copy the text in the dialog box that opens.
2. Paste the text into the support web form.

Glossary

Attachment—Combination of an attachment position and a reagent file.

Attachment position—Position on a core to which side chains are attached (synonymous with ‘site of diversity’).

Chemical feature—One of six standard types of functional group: hydrogen bond acceptor (A), hydrogen bond donor (D), hydrophobe (H), negatively charged center (N), positively charged center (P), or aromatic (R). The CombiGlide definition of these feature types differs somewhat from that of standard pharmacophore analysis.

Combinatorial library—A set of compounds prepared by combinatorial synthesis.

Combinatorial synthesis—The synthesis of sets of compounds utilizing the same reaction sequence where all possible combinations of the reagents are used.

Compound—In CombiGlide, a compound is explicitly defined by a combination of reagents that are attached to a core to form a product. A compound may encompass several structures. If a reagent exists in more than one form (different tautomers, conformers, stereoisomers, protonation state), the compound may also encompass more than one variant.

Core—The structural element of a combinatorial library that is constant throughout the library. It is the part of the core-containing molecule that remains when the side chains are removed.

Enumerate—To generate structures of all possible members of a combinatorial library.

Library selection—The results of filtering and selection that generates a list of reagents at each position for a focused library, including information on the filtering and selection strategy.

Fully substituted core—A core with side chains attached at all attachment positions.

Library definition—A collection of the core-containing structure and the reagent files that defines a combinatorial library. Can be used as input to library enumeration or combinatorial screening.

Minimal capping group—The smallest representative side chain for a particular functional group.

Minimally capped core—A core with the appropriate minimal capping group at each attachment position.

Protocore—A molecule that could be used to replace the core molecule, and that has hydrogen atoms where the side chains should be attached.

Reagent—The chemical reagent used in the combinatorial synthesis; a structure or set of structures with the same chemical composition and heavy-atom connectivity that is used to add a side chain to the core. A reagent may encompass more than one structure if it exists in more than one form (different tautomers, conformers, stereoisomers, protonation state).

Reagent variant—A single reagent structure, defined by a particular tautomer, conformer, stereoisomer, and protonation state of the reagent.

Side chains—The structural elements of a combinatorial library that are varied within the library. The side chain is the part of a reagent that remains after removing the reactive functional group that is discarded when it is bound to the core.

Structure—A molecule with a particular geometry and stereochemistry, in a single tautomeric form, ionization state, and conformation.

User-capped core—The molecule that was used to define the core. The capping groups are the side chains from this molecule.

Type—The list of chemical features that can be found in a side chain.

A

actives	
enrichment data	81
specifying	80
atoms, picking for grow bond	98
attachment	39, 44
defining	19, 45
deleting	45
attachment position	39, 44
defining	19, 45
restrictions on	22

B

Boltzmann population	62
----------------------------	----

C

center of mass, constraining core	53
chemical features	
adding to reagent file	101
command-line analysis	102
definition	74
frequencies of occurrence	75
Collections panel	12
CombiGlide XP docking	3
combinatorial docking	40, 49
combinatorial library	47
docking	57
enumerating	17
compounds	
definition	109
number in library	66
sources of	22
conventions, document	vii
core poses	
constraints	53
selecting molecules for	54
use in docking	55
core, definition	39
core-containing molecules	
criteria for core poses	54
preparing	36
Create Fragment Collection panel	15

D

directory	
installation	5
Maestro working	5
docking	
combinatorial	40, 49
exporting results	59
maximum number of structures	57
options for	50
single-position	40, 49

E

enrichment	
of actives	81
of chemical features	75
environment variable	
SCHRODINGER	5

F

filters	
custom	66
preset	65
fragment names, generating list of	89
functional groups	
choosing	28
file format	96
multiple instances	27
storage location	95
table	29

G

Glide docking modes	57
grow bond	39, 44, 95
grow names	96

H

hydrophobe, definition	74
------------------------------	----

I

Interactive Enumeration panel	
Combinatorial Library tab	10
Fragment Collection tab	13
ionization states, generating for reagents	34

L

library definition, saving	47
library selections	66
saving	68
library size	66
library, designing	79
limits, docking	50, 57
log file, reagent preparation	35
long name, reagent	95

M

metal-binding states, generating for reagents ...	34
minimal capping group	
definition	40
modifying	99
structure defining	95
use in docking	49
minimally capped core	40, 45

O

output files, reagent preparation	35
---	----

P

path lengths	102
product installation	106
project, saving	37
protein structures, obtaining and importing	22

Q

QikProp properties	
filtering on	65
use in hydrophobe definition	74

R

reagent files	
generating	25
merging	103
selecting	19, 45
reagent type	26
reagents	
generating ionization, tautomeric, and metal-binding states	34
long name	95
minimum and maximum selected	63

prefiltering	22
restrictions on structures	27
short name	95
sources of	22
titles	27

receptor grid, selecting	50
--------------------------------	----

restrictions	
attachment positions	22
in docking	50, 57
reagent structures	27
run, definition	37

S

Schrödinger contact information	106
selection strategies	
Boltzmann population	62
GlideScore	62
single-position docking	61
short name, reagent	95
side chains	
chemical feature type	74
diversity of	75
single-position docking	40, 49
SMARTS patterns, for functional group definition	95
structures	
2D-to-3D conversion	34
core-containing	43
number in library	66
reagent, varying	34
restrictions on	27
tangled	70

T

tautomers, generating for reagents	34
--	----

U

untangling structures	19, 70
user-capped core	54

120 West 45th Street
17th Floor
New York, NY 10036

155 Gibbs St
Suite 430
Rockville, MD 20850-0353

Quatro House
Frimley Road
Camberley GU16 7ER
United Kingdom

101 SW Main Street
Suite 1300
Portland, OR 97204

Dynamostraße 13
D-68165 Mannheim
Germany

8F Pacific Century Place
1-11-1 Marunouchi
Chiyoda-ku, Tokyo 100-6208
Japan

245 First Street
Riverview II, 18th Floor
Cambridge, MA 02142

Zeppelinstraße 73
D-81669 München
Germany

No. 102, 4th Block
3rd Main Road, 3rd Stage
Sharada Colony
Basaveshwaranagar
Bangalore 560079, India

8910 University Center Lane
Suite 270
San Diego, CA 92122

Potsdamer Platz 11
D-10785 Berlin
Germany

SCHRÖDINGER®